

Operating Systems (CSE531)

Lecture # 08



Manish Shrivastava

LTRC, IIIT Hyderabad

Outline

- Cooperating Processes
- Inter-process Communication
 - Shared Memory
 - Message Passing
- Communication between separate processes
- Methods to implement a link
- Client Server Communication
 - Sockets
 - Remote Procedure Call (RPC)
 - Pipes

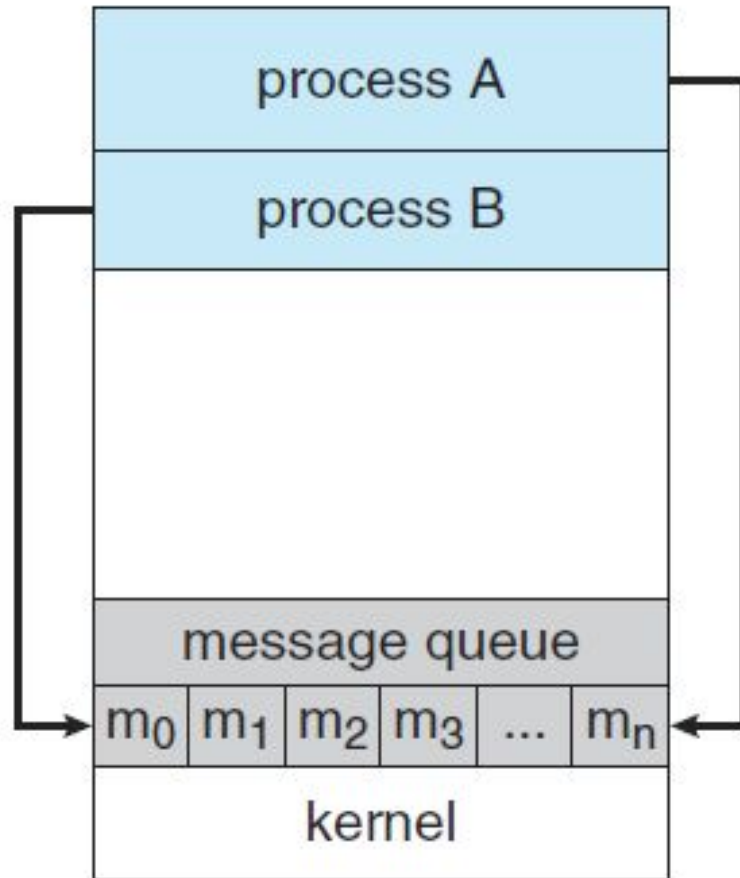
Cooperating Processes

- The processes can be independent or cooperating processes.
- *Independent* process **cannot** affect or be affected by the execution of another process.
- *Cooperating* process **can** affect or be affected by the execution of another process
- Advantages of process cooperation
 - Information sharing
 - Computation speed-up: Break into several subtasks and run in parallel
 - Modularity: Constructing the system in modular fashion.
 - Convenience: User will have many tasks to work simultaneously
 - Editing, compiling, printing

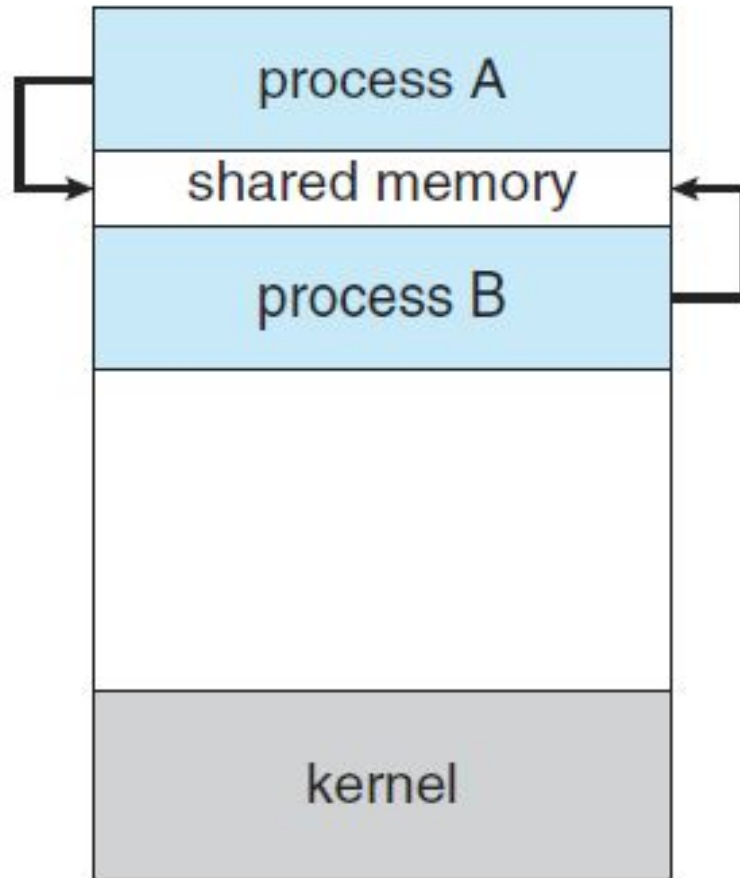
Inter-process Communication (IPC)

- IPC facility provides a mechanism to allow processes to communicate and synchronize their actions.
- Processes can communicate through
 - **Shared Memory**
 - **Message Passing.**Both schemes may exist in OS.
- The Shared-memory method requires communication processes to share some variables.
 - The responsibility for providing communication rests with the programmer.
 - The OS only provides shared memory.
- Example: producer-consumer problem.

Communication Models



Message Passing



Shared Memory

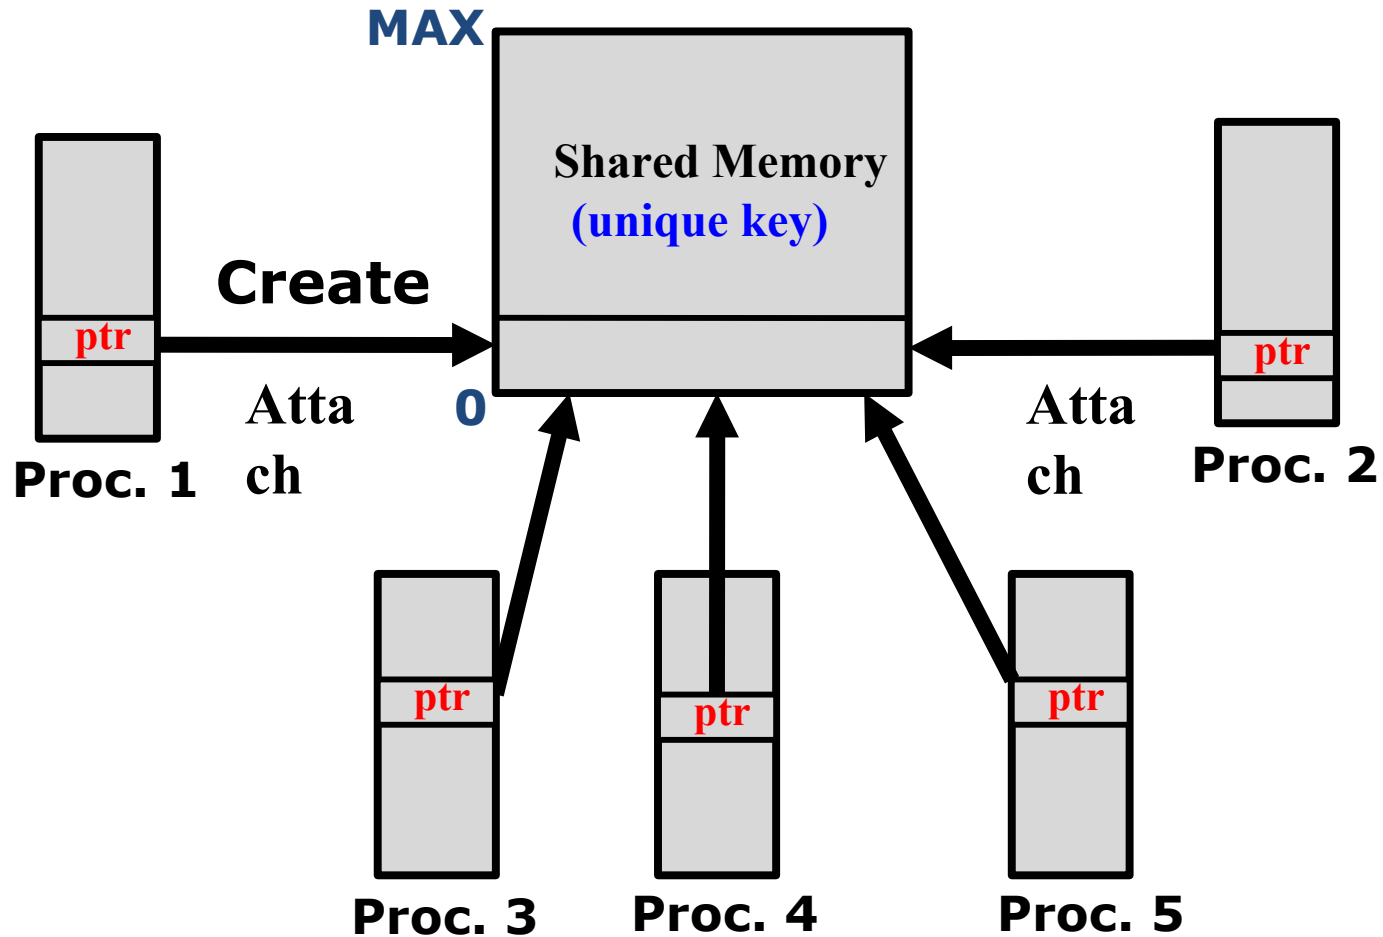
Producer-Consumer Problem: Shared memory

- *Producer* process produces information that is consumed by a *consumer* process.
 - *unbounded-buffer* places no practical limit on the size of the buffer.
 - Producer can produce any number of items.
 - Consumer may have to wait
 - *bounded-buffer* assumes that there is a fixed buffer size.
 - Consumer or producer may have to wait

Shared Memory

- Shared memory is memory that may be simultaneously accessed by multiple programs.
- Life cycle of Shared Memory
 - Create Memory Block
 - Attach to number of process simultaneously
 - Read & Write
 - Detach Memory block
 - Free Memory block

Common chunk of read/write memory



Example

Shmget System Call

The shmget system call is used to create shared memory segment

Shmat System Call

Attaches block to process

```
int shmget (key_t key, int size, int shmflg);
```

Arguments:

- *key_t key*: key for creating or accessing shared memory
- *int size*: size in bytes of shared memory segment to create. Use 0 for accessing an existing segment.
- *int shmflg*: segment creation condition and access permission.

Communicating Among Separate Processes

Define the structure of a shared memory segment as follows:

```
//Status of memory block
#define NOT_READY (-1)
#define FILLED (0)
#define TAKEN (1)

struct Memory {
    int status; // used for Synchronization
    int data[4]; // data to share
};
```

Communicating Among Separate Processes

The “Server”

*Prepare for a shared
memory*



```
void main(int argc, char *argv[])
{
    key_t      ShmKEY;
    int        ShmID, i;
    struct Memory *ShmPTR;

    ShmKEY = ftok("./", 'x');
    ShmID = shmget(ShmKEY, sizeof(struct Memory),
                  IPC_CREAT | 0666);
    ShmPTR = (struct Memory *) shmat(ShmID, NULL, 0);
```

Communicating Among Separate Processes

shared memory not ready



```
ShmPTR-->status = NOT_READY;
```

filling in data

```
for (i = 0; i < 4; i++)  
    ShmPTR- -> data[i] = atoi(argv[i]);
```

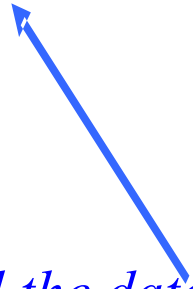
```
ShmPTR-->status = FILLED;  
while (ShmPTR-->status != TAKEN)  
    sleep(1);  /* sleep for 1 second */
```

```
shmdt((void *) ShmPTR);  
shmctl(ShmID, IPC_RMID, NULL);  
exit(0);
```

detach and remove shared memory



wait until the data is taken



Communicating Among Separate Processes

```
void main(void)
```

```
{
```

```
    key_t          ShmKEY;
```

```
    int            ShmID;
```

```
    struct Memory  *ShmPTR;
```

The “Client”

*prepare for shared
memory*



```
    ShmKEY=ftok(".", 'x');
```

```
    ShmID = shmget(ShmKEY, sizeof(struct Memory), 0666);
```

```
    ShmPTR = (struct Memory *) shmat(ShmID, NULL, 0);
```

```
    while (ShmPTR->status != FILLED)
```

```
    ;
```

```
    printf("%d %d %d %d\n", ShmPTR-->data[0],
```

```
           ShmPTR-->data[1], ShmPTR-->data[2], ShmPTR-->data[3]);
```

```
    ShmPTR->status = TAKEN;
```

```
    shmdt((void *) ShmPTR);
```

```
    exit(0);
```

```
}
```

Communicating Among Separate Processes

- If you did not remove your shared memory segments (*e.g.*, program crashes before the execution of `shmctl()`), they will be in the system forever. This will degrade the system performance.
- Use the `IpCs` command to check if you have shared memory segments left in the system.
- Use the `ipcrm` command to remove your shared memory segments.

Inter-process Communication (IPC): Message Passing System

- **Message system** – processes communicate with each other without resorting to shared variables.
- If P and Q want to communicate, a communication link exists between them.
- OS provides this facility.

IPC

- IPC facility provides two operations:
 - **send**(*message*) – message size fixed or variable
 - **receive**(*message*)
- If P and Q wish to communicate, they need to:
 - establish a *communication link* between them
 - exchange messages via send/receive
- Implementation of communication link
 - physical (e.g., shared memory, hardware bus)
 - logical (e.g., logical properties)
 - We are concerned with logical link.

Implementation Questions

- How are links established?
- Can a link be associated with more than two processes?
- How many links can there be between every pair of communicating processes?
- What is the capacity of a link?
- Is a link unidirectional or bi-directional?
- Is the size of a message that the link can accommodate fixed or variable?

Fixed and variable message size

- Fixed size message
 - Good for OS designer
 - Complex for programmer
- Variable size messages
 - Complex for the OS designer
 - Good for programmer

Methods to implement a link

- Direct or Indirect communication
- Synchronous or asynchronous communication
- Automatic or explicit buffering

Direct Communication

- **Symmetric:** Processes must name each other explicitly:
 - `send (P, message)` – send a message to process P
 - `receive(Q, message)` – receive a message from process Q.
- **Asymmetric:** Only sender names the recipient, the recipient is not required to name the sender.
 - The send and receive primitives are as follows.
 - `Send (P, message)`– send a message to process P.
 - `Receive(id, message)`– receive a message from any process.
- Properties of communication link
 - Links are established automatically.
 - A link is associated with exactly one pair of communicating processes.
 - Between each pair there exists exactly one link.
 - The link may be unidirectional, but is usually bi-directional
- **Disadvantages:** Changing the identifier of a process may necessitate examining all other process definitions. Any such **hard-coding** techniques, where identifiers must be explicitly stated, are less desirable.

Indirect Communication

- The messages are sent and received from *mailboxes* (also referred to as *ports*).
- A mailbox is an object
 - Process can place messages
 - Process can remove messages.
- Two processes can communicate only if they have a shared mailbox.
- Operations
 - create a new mailbox
 - send and receive messages through mailbox
 - destroy a mailbox
- Primitives are defined as:
 - send**(*A, message*) – send a message to mailbox A
 - receive**(*A, message*) – receive a message from mailbox A

Indirect Communication

- Mailbox sharing
 - P_1 , P_2 , and P_3 share mailbox A.
 - P_1 , sends; P_2 and P_3 receive.
 - Who gets a message ?
- Solutions
 - Allow a link to be associated with at most two processes.
 - Allow only one process at a time to execute a receive operation.
 - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.
- A mailbox may be owned either by a process or by the operating system.
 - If the mailbox is part of the address space of the process,
 - We distinguish between the owner (which can only receive messages through this mailbox) and the user (which can only send messages to the mailbox).
 - Since each mailbox has a unique owner, there can be no confusion about which process should receive a message sent to this mailbox.
 - When a process that owns a mailbox terminates, the mailbox disappears.

Indirect Communication

- Properties of a link:
 - A link is established if they have a shared mailbox
 - A link may be associated with more than two boxes.
 - Between a pair of processes there may be number of links
 - A link may be either unidirectional or bi-directional.
- OS owned mailboxes:
 - OS provides a facility
 - To create a mailbox
 - Send and receive messages through mailbox
 - To destroy a mail box.
- The process that creates mailbox is a owner of that mailbox
- The ownership and send and receive privileges can be passed to other processes through system calls.

Methods to implement a link

- Direct or Indirect communication
 - Direct link vs Mailbox
- Synchronous or asynchronous communication
- Automatic or explicit buffering

Synchronous v/s asynchronous

- Message passing may be either blocking or non-blocking.
- **Blocking** is considered **synchronous**
- **Non-blocking** is considered **asynchronous**
- **send** and **receive** primitives may be either blocking or non-blocking.
 - **Blocking send:** The sending process is blocked until the message is received by the receiving process or by the mailbox.
 - **Non-blocking send:** The sending process sends the message and resumes operation.
 - **Blocking receive:** The receiver blocks until a message is available.
 - **Non-blocking receive:** The receiver receives either a valid message or a null.

Methods to implement a link

- Direct or Indirect communication
 - Direct link vs Mailbox
- Synchronous or asynchronous communication
- Automatic or explicit buffering

Automatic and explicit buffering

- A link has some capacity that determines the number of messages that can reside in it temporarily.
- Queue of messages is attached to the link; implemented in one of three ways.
 1. **Zero capacity** – 0 messages
Sender must wait for receiver (rendezvous).
 2. **Bounded capacity** – finite length of n messages
Sender must wait if link full.
 3. **Unbounded capacity** – infinite length
Sender never waits.
- In non-zero capacity cases a process does not know whether a message has arrived after the send operation.

Automatic and explicit buffering

- The sender must communicate explicitly with receiver to find out whether the later received the message.
- Example: Suppose P sends a message to Q and executes only after the message has arrived.
- Process P:
 - `send (Q, message)` : send message to process Q
 - `receive(Q,message)` : Receive message from process Q
- Process Q
 - `Receive(P,message)`
 - `Send(P,"ack")`

Exception conditions

- When a failure occurs error recovery (exception handling) must take place.
- Process termination
 - A sender or receiver process may terminate before a message is processed. (may be blocked forever)
 - A system will terminate the other process or notify it.
- Lost messages
 - Messages may be lost over a network
 - Timeouts; restarts.
- Scrambled messages
 - Message may be scrambled on the way due to noise
 - The OS will retransmit the message
 - Error-checking codes (parity check) are used.

Client-Server Communication

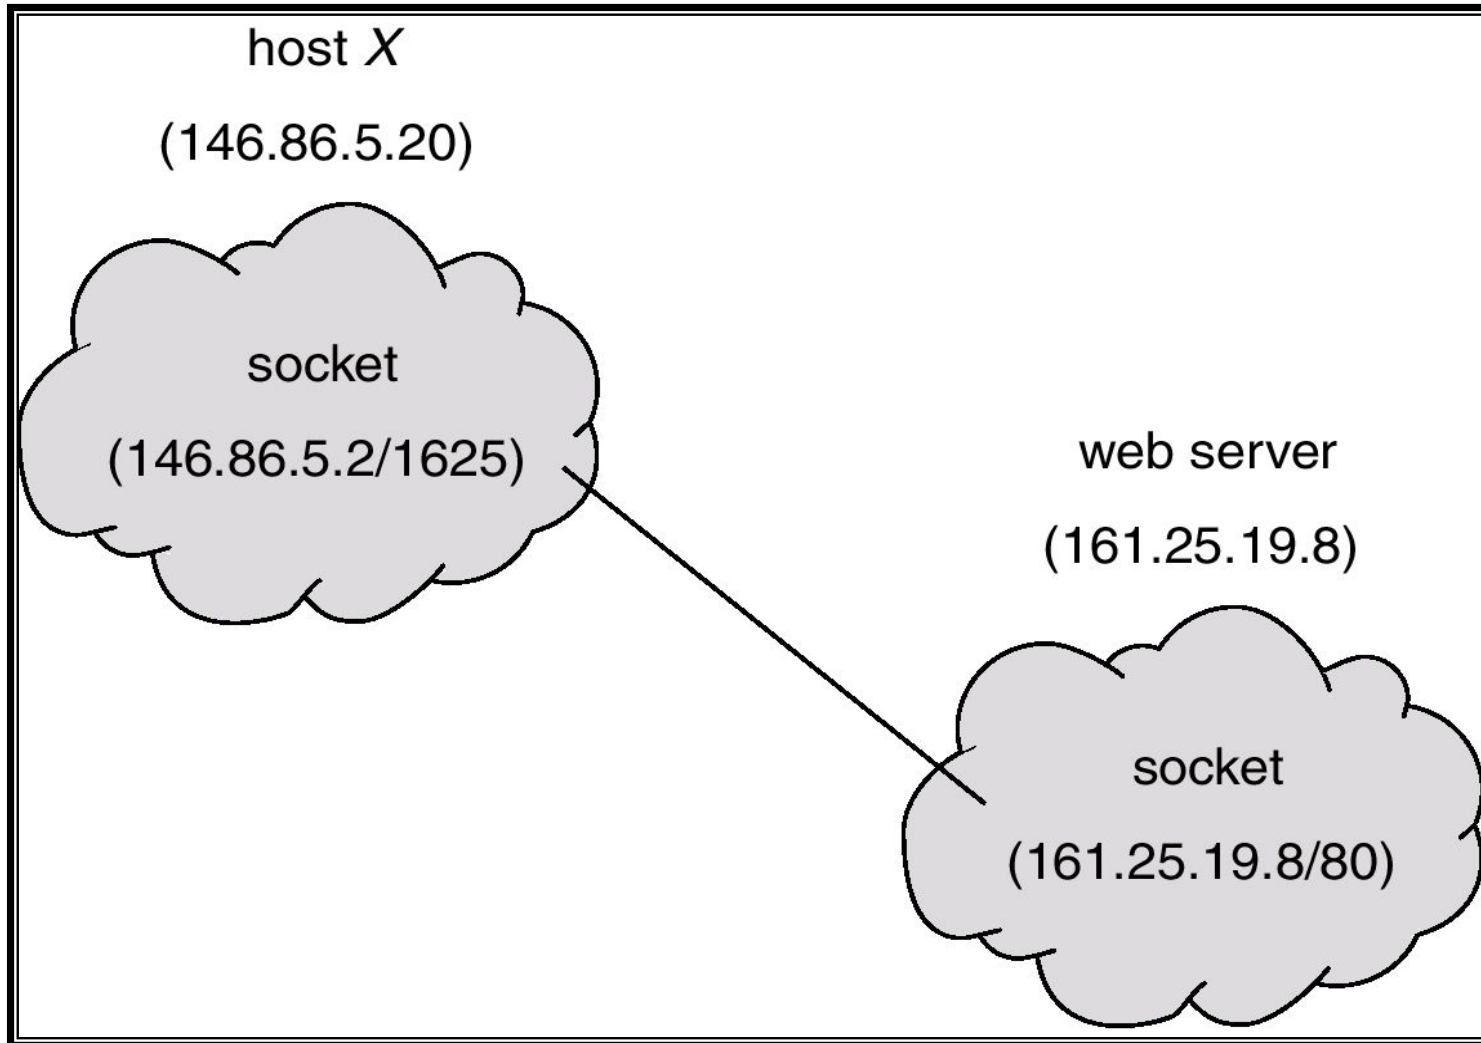
Strategies for communication in client–server systems

- Sockets
- Remote Procedure Calls
- Pipes

Sockets

- A socket is defined as an endpoint for communication.
- A pair of processes communicating over a network employs **a pair of sockets**— one for each process.
- Socket: ***Concatenation of IP address and port***
- The socket 161.25.19.8:1625 refers to port 1625 on host 161.25.19.8
- Servers implementing specific services listen to well-known ports
 - telnet server listens to port 23
 - ftp server listens to port 21
 - http server listens to port 80.
- The ports less than 1024 are used for standard services.
- The port for socket is an arbitrary number greater than 1024.
- Communication happens between a pair of sockets.

Sockets



Socket: Example Code

- Server

```
import java.net.*;
import java.io.*;

public class DateServer
{
    public static void main(String[] args) {
        try {
            ServerSocket sock = new ServerSocket(6013);

            /* now listen for connections */
            while (true) {
                Socket client = sock.accept();

                PrintWriter pout = new
                    PrintWriter(client.getOutputStream(), true);

                /* write the Date to the socket */
                pout.println(new java.util.Date().toString());

                /* close the socket and resume */
                /* listening for connections */
                client.close();
            }
        } catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

- Client

```
import java.net.*;
import java.io.*;

public class DateClient
{
    public static void main(String[] args) {
        try {
            /* make connection to server socket */
            Socket sock = new Socket("127.0.0.1", 6013);

            InputStream in = sock.getInputStream();
            BufferedReader bin = new
                BufferedReader(new InputStreamReader(in));

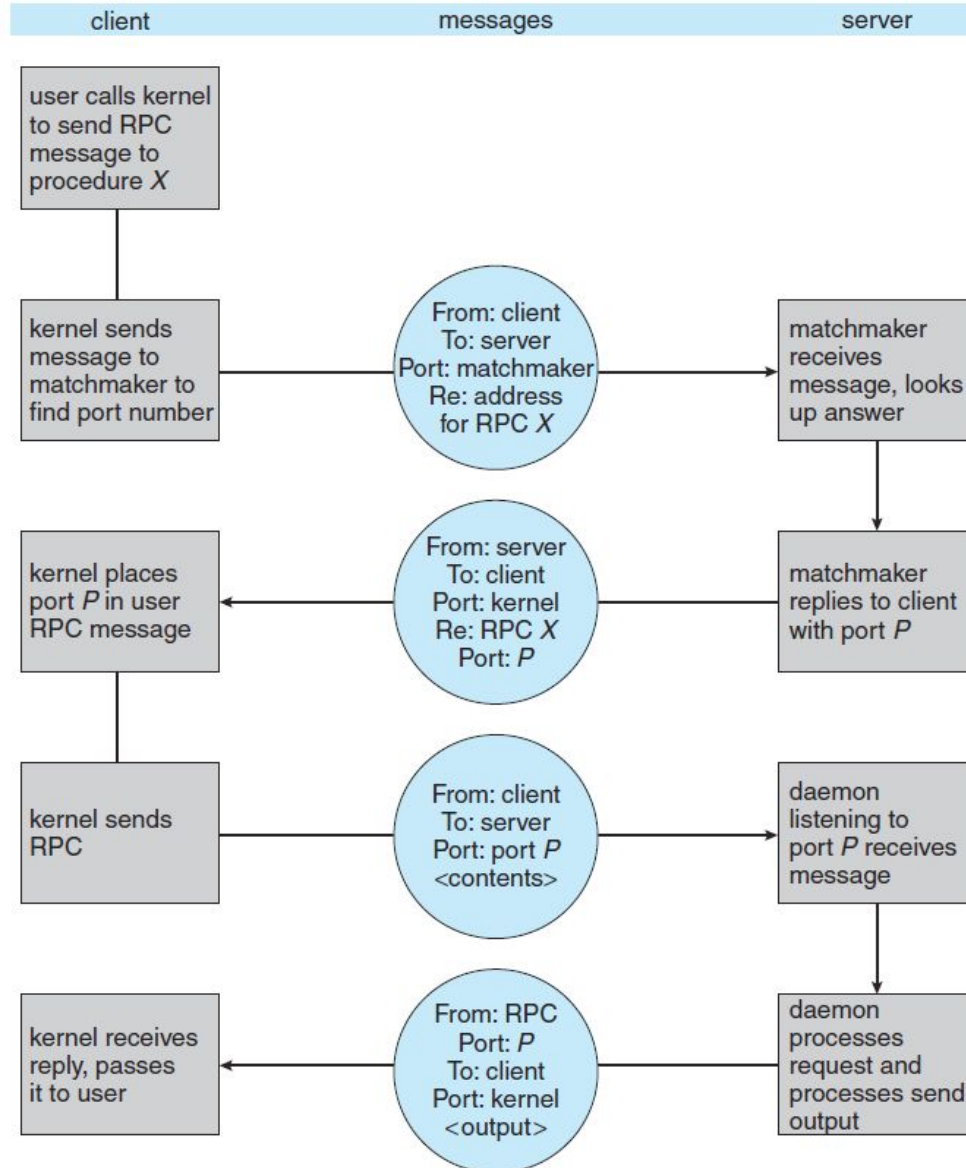
            /* read the date from the socket */
            String line;
            while ( (line = bin.readLine()) != null)
                System.out.println(line);

            /* close the socket connection */
            sock.close();
        } catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

Remote Procedure Calls

- Remote procedure call (RPC) abstracts procedure calls between processes on networked systems.
- The semantics of RPCs allows a client to invoke a procedure on a remote host as it would invoke a procedure locally
- **Stubs** – client-side proxy for the actual procedure on the server.
- The client-side stub locates the server and *marshals* the parameters.
 - Marshalling: Parameter marshalling involves packaging the parameters into a form that can be transmitted over a network.
- The server-side stub receives this message, unpacks the marshaled parameters, and performs the procedure on the server.

Execution of RPC



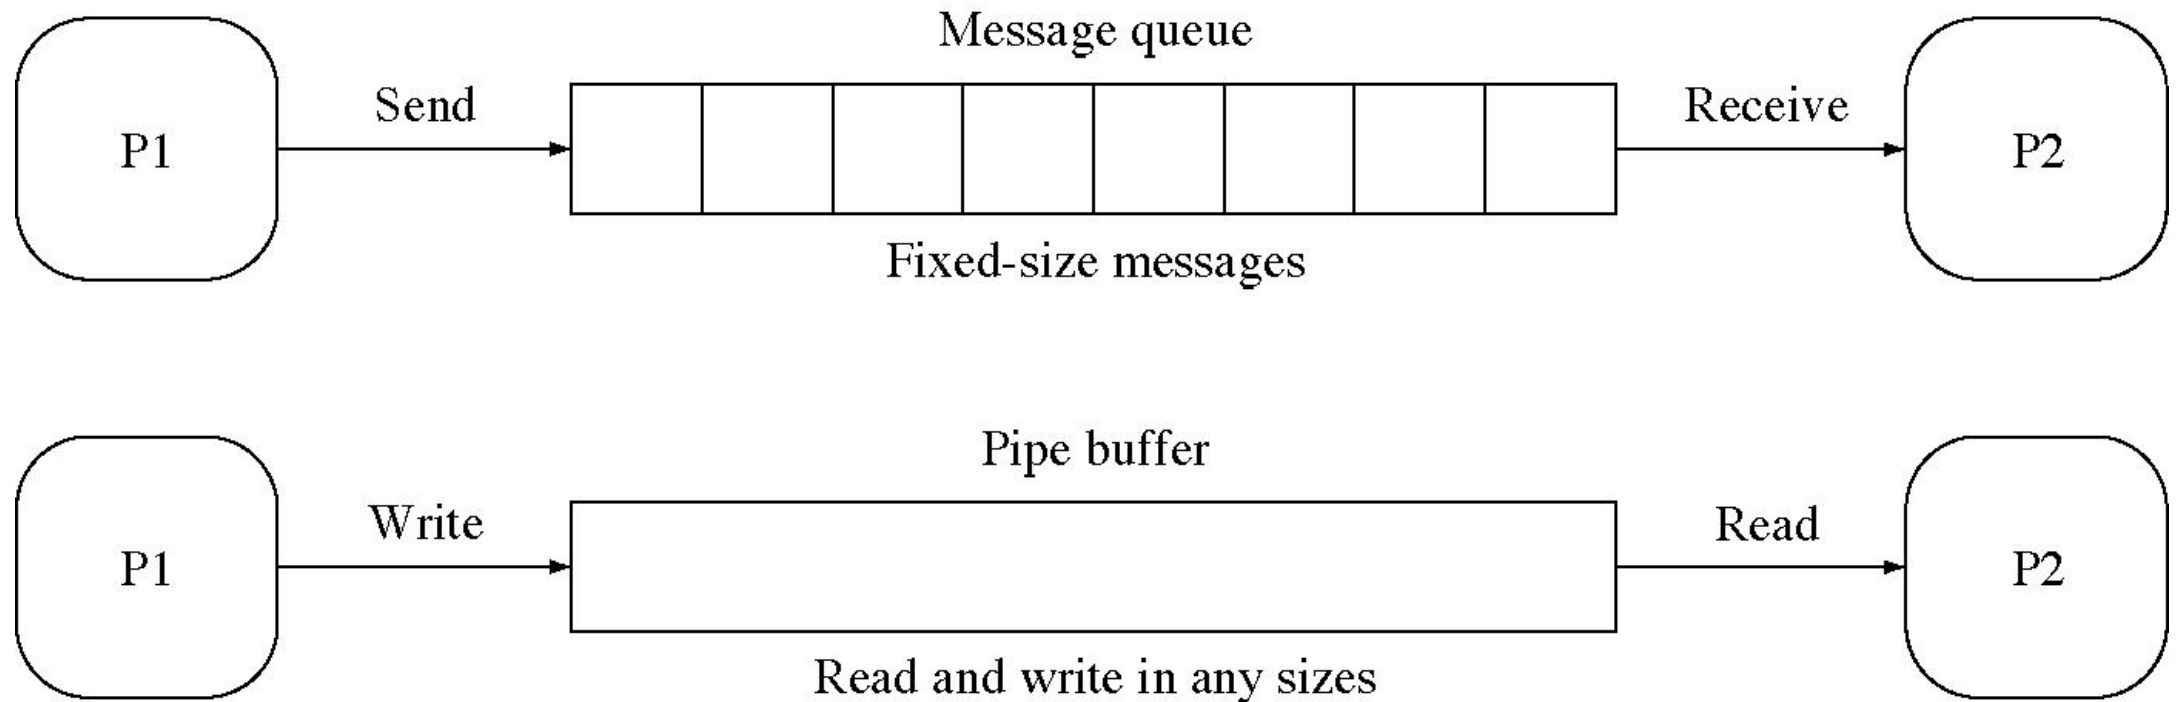
Remote Procedure Calls

- There are several issues
 - Whereas local procedure calls fail only under extreme circumstances, RPCs can fail, or be duplicated and executed more than once, as a result of common network errors.
 - Solution
 - In RPC, the server must implement “at most once” semantics send the ack to client to ensure no further communication.
 - The server detects the repeated messages.
 - Even client sends more messages but those will be ignored.
 - Binding issues: linking, loading and execution
 - Fixed port addresses or rendezvous
- Applications: distributed file systems

Pipes

- Pipe: one of the first IPC mechanisms in early UNIX systems
 - uses the familiar file interface
 - not a special interface (like messages)
- Connects an open file of one process to an open file of another process
 - Often used to connect the standard output of one process to the standard input of another process

Messages v/s Pipes




More about Pipes

- Acts as a conduit allowing two processes to communicate
- **Issues**
 - Is communication unidirectional or bidirectional?
 - In the case of two-way communication, is it half-duplex (data can travel only one way at a time) or full-duplex?
 - Must there exist a relationship (i.e. parent-child) between the communicating processes?
 - Can the pipes be used over a network?

Ordinary Pipes

- **Ordinary Pipes** allow communication in standard producer-consumer style
- Producer writes to one end (the *write-end* of the pipe)
- Consumer reads from the other end (the *read-end* of the pipe)
- Ordinary pipes are therefore unidirectional
- Require parent-child relationship between communicating processes
 - A parent process creates a pipe and uses it to communicate with a child process that it creates via `fork()`.
- Ordinary pipe can not be accessed from outside the process that creates it.
- Once the processes have finished communicating and have terminated, the ordinary pipe ceases to exist.

Ordinary Pipes



- Unix function to create pipes: `pipe(int fd[])`
- `fd[0]` is the read end. `fd[1]` is write end.

Example Code

```
#include <sys/types.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>

#define BUFFER_SIZE 25
#define READ_END 0
#define WRITE_END 1

int main(void)
{
    char write_msg[BUFFER_SIZE] = "Greetings";
    char read_msg[BUFFER_SIZE];
    int fd[2];
    pid_t pid;

    /* create the pipe */
    if (pipe(fd) == -1) {
        fprintf(stderr, "Pipe failed");
        return 1;
    }

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
}
```

```
if (pid > 0) { /* parent process */
    /* close the unused end of the pipe */
    close(fd[READ_END]);

    /* write to the pipe */
    write(fd[WRITE_END], write_msg, strlen(write_msg)+1);

    /* close the write end of the pipe */
    close(fd[WRITE_END]);
}
else { /* child process */
    /* close the unused end of the pipe */
    close(fd[WRITE_END]);

    /* read from the pipe */
    read(fd[READ_END], read_msg, BUFFER_SIZE);
    printf("read %s", read_msg);

    /* close the write end of the pipe */
    close(fd[READ_END]);
}

return 0;
}
```

Named Pipes

- Named Pipes are more powerful than ordinary pipes
- Communication is bi-directional
- No parent-child relationship is necessary between the communicating processes
- Several processes can use the named pipe for communication
 - Typically, a named pipe has several writers
- Named pipes continue to exist after communicating processes have finished.
- Named pipes are referred to as FIFOs in UNIX systems
 - Only half-duplex transmission is permitted.
 - The communicating processes must reside on the same machine, else one should use sockets.

Reference

- **Beej's Guide to Unix IPC:**

<http://beej.us/guide/bgipc/output/html/multipage/index.html>

THANK YOU