


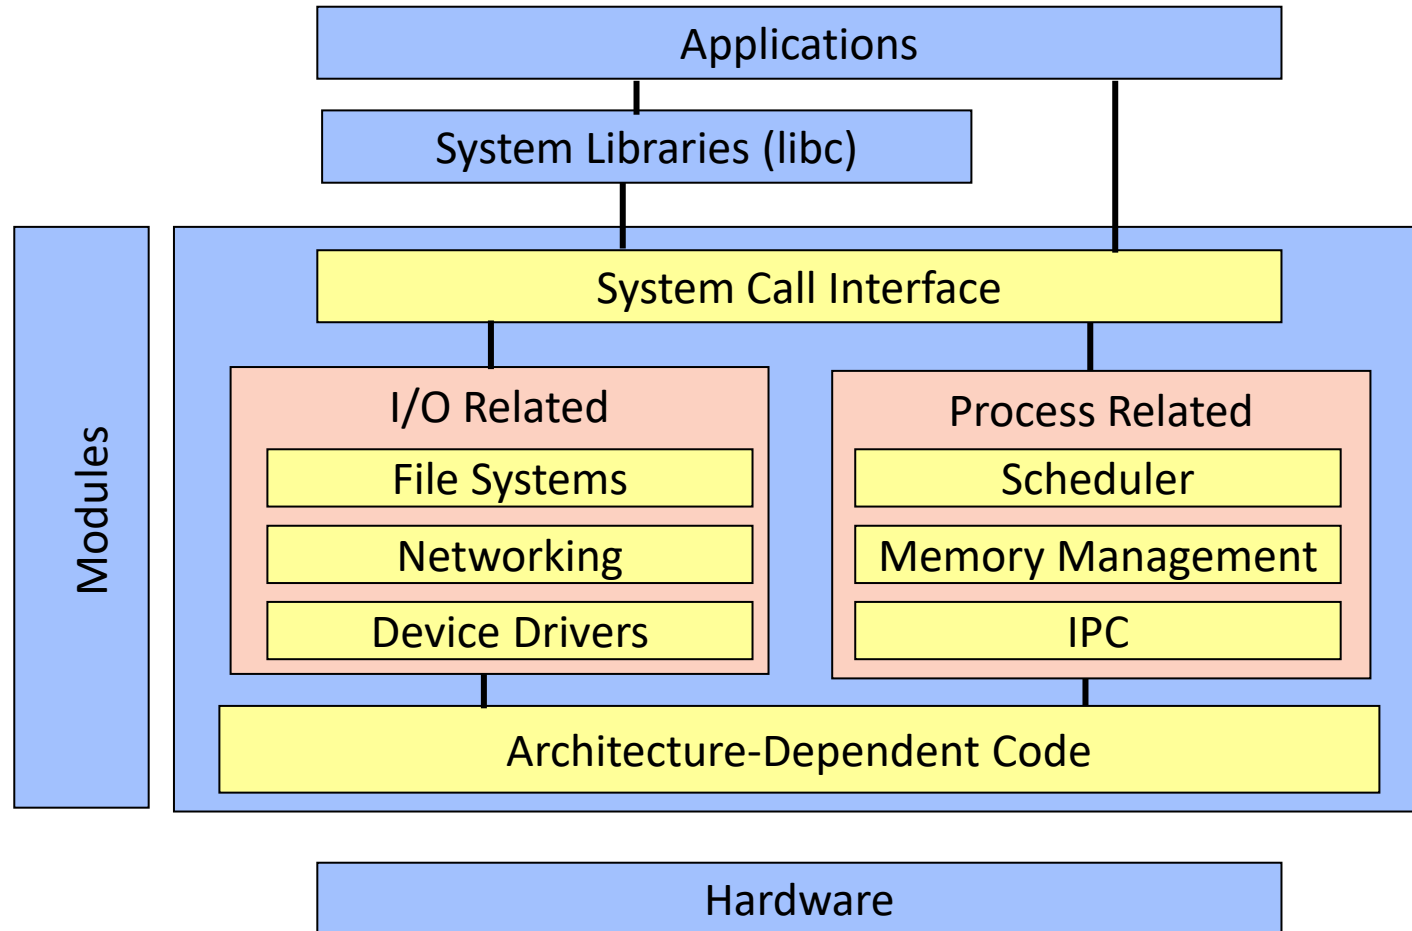
The Linux Kernel: Introduction



Kernel Design Goals

- Performance: efficiency, speed.
 - Utilize resources to capacity with low overhead.
- Stability: robustness, resilience.
 - Uptime, graceful degradation.
- Capability: features, flexibility, compatibility.
- Security, protection.
 - Protect users from each other & system from bad users.
- Portability.
- Extensibility.

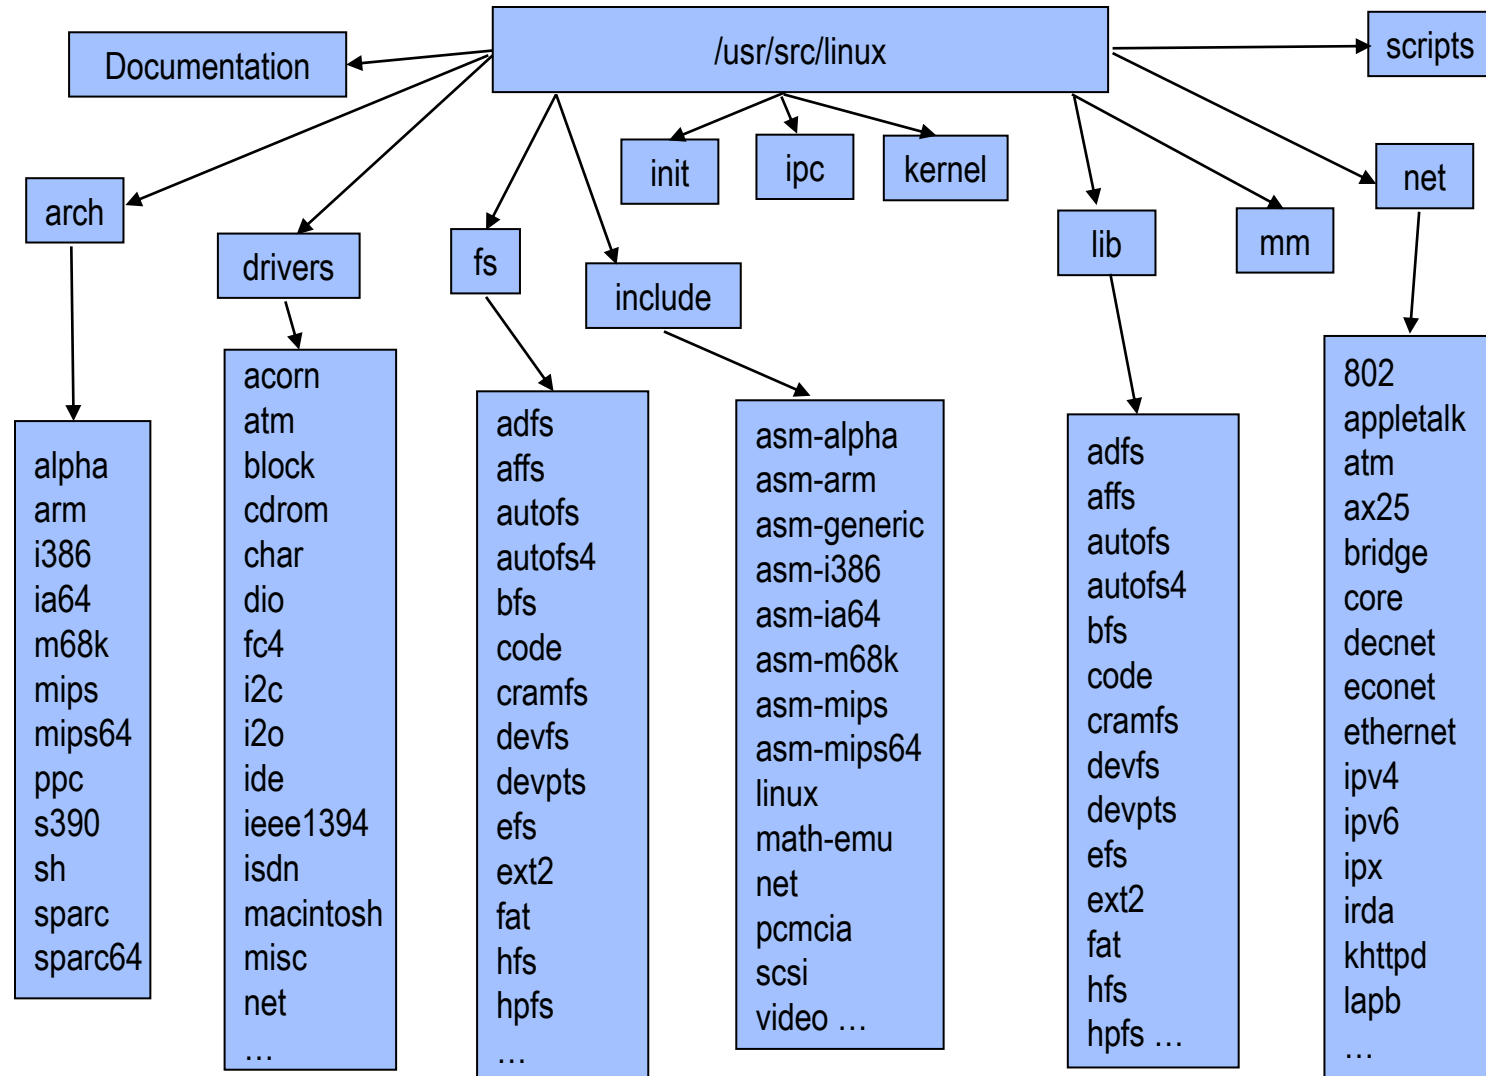
Example “Core” Kernel



Architectural Approaches

- Monolithic.
- Layered.
- Modularized.
- Micro-kernel.
- Virtual machine.

Linux Source Tree Layout



linux/arch

- Subdirectories for each current port.
- Each contains **kernel**, **lib**, **mm**, **boot** and other directories whose contents override code stubs in architecture independent code.
- **lib** contains highly-optimized common utility routines such as memcpy, checksums, etc.
- **arch** as of 2.4:
 - alpha, arm, i386, ia64, m68k, mips, mips64.
 - ppc, s390, sh, sparc, sparc64.

linux/drivers

- Largest amount of code in the kernel tree (~1.5M).
- device, bus, platform and general directories.
- drivers/char – n_tty.c is the default line discipline.
- drivers/block – elevator.c, genhd.c, linear.c, ll_rw_blk.c, raidN.c.
- drivers/net –specific drivers and general routines Space.c and net_init.c.
- drivers/scsi – scsi_*.c files are generic; sd.c (disk), sr.c (CD-ROM), st.c (tape), sg.c (generic).
- General:
 - cdrom, ide, isdn, parport, pcmcia, pnp, sound, telephony, video.
- Buses – fc4, i2c, nubus, pci, sbus, tc, usb.
- Platforms – acorn, macintosh, s390, sgi.

linux/fs

- Contains:
 - virtual filesystem (VFS) framework.
 - subdirectories for actual filesystems.
- vfs-related files:
 - exec.c, binfmt_*.c - files for mapping new process images.
 - devices.c, blk_dev.c – device registration, block device support.
 - super.c, filesystems.c.
 - inode.c, dcache.c, namei.c, buffer.c, file_table.c.
 - open.c, read_write.c, select.c, pipe.c, fifo.c.
 - fcntl.c, ioctl.c, locks.c, dquot.c, stat.c.

linux/include

- include/asm-*:
 - Architecture-dependent include subdirectories.
- include/linux:
 - Header info needed both by the kernel and user apps.
 - Usually linked to /usr/include/linux.
 - Kernel-only portions guarded by #ifdefs
 - #ifdef __KERNEL__
 - /* kernel stuff */
 - #endif
- Other directories:
 - math-emu, net, pcmcia, scsi, video.

linux/init

- Just two files: version.c, main.c.
- version.c – contains the version banner that prints at boot.
- main.c – architecture-independent boot code.
- start_kernel is the primary entry point.

linux/ipc

- System V IPC facilities.
- If disabled at compile-time, util.c exports stubs that simply return – ENOSYS.
- One file for each facility:
 - sem.c – semaphores.
 - shm.c – shared memory.
 - msg.c – message queues.

linux/kernel

- The core kernel code.
- sched.c – “the main kernel file”:
 - scheduler, wait queues, timers, alarms, task queues.
- Process control:
 - fork.c, exec.c, signal.c, exit.c etc...
- Kernel module support:
 - kmod.c, ksyms.c, module.c.
- Other operations:
 - time.c, resource.c, dma.c, softirq.c, itimer.c.
 - printk.c, info.c, panic.c, sysctl.c, sys.c.

linux/lib

- kernel code cannot call standard C library routines.
- Files:
 - brlock.c – “Big Reader” spinlocks.
 - cmdline.c – kernel command line parsing routines.
 - errno.c – global definition of errno.
 - inflate.c – “gunzip” part of gzip.c used during boot.
 - string.c – portable string code.
 - Usually replaced by optimized, architecture-dependent routines.
 - vsprintf.c – libc replacement.

linux/mm

- Paging and swapping:
 - swap.c, swapfile.c (paging devices), swap_state.c (cache).
 - vmscan.c – paging policies, kswapd.
 - page_io.c – low-level page transfer.
- Allocation and deallocation:
 - slab.c – slab allocator.
 - page_alloc.c – page-based allocator.
 - vmalloc.c – kernel virtual-memory allocator.
- Memory mapping:
 - memory.c – paging, fault-handling, page table code.
 - filemap.c – file mapping.
 - mmap.c, mremap.c, mlock.c, mprotect.c.

linux/scripts

- Scripts for:
 - Menu-based kernel configuration.
 - Kernel patching.
 - Generating kernel documentation.


Summary

- Linux is a modular, UNIX-like monolithic kernel.
- Kernel is the heart of the OS that executes with special hardware permission (kernel mode).
- “Core kernel” provides framework, data structures, support for drivers, modules, subsystems.
- Architecture dependent source sub-trees live in /arch.

The Bootloader

The Bootloader

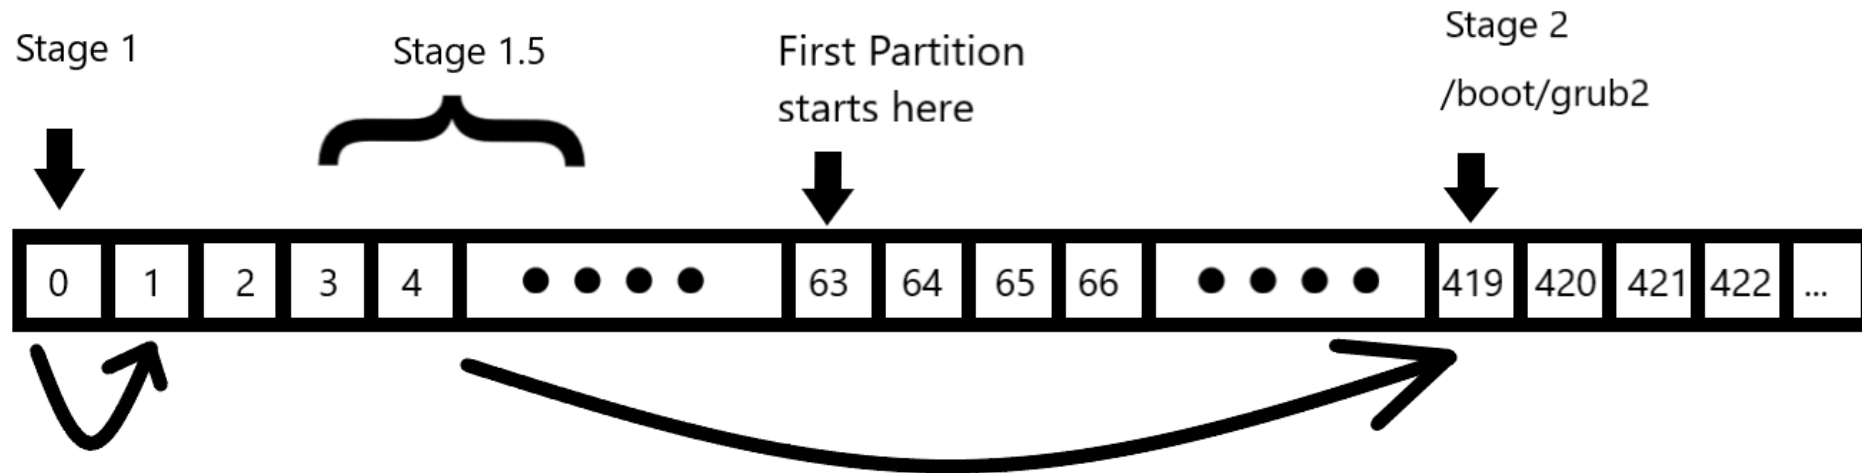
- Responsible for loading the operating system into memory.
- LILO (Linux Loader) was the default for a long time
- It works by pointing to the first sector where the kernel is stored (simplified).



- Has limited functionality due to the 446 bytes available in MBR

The Bootloader - GRUB

- GRUB replaced LILO as the default Linux bootloader (we'll focus on version 2 but version 1 is similar).
- Split into 3 stages.



The Bootloader - GRUB

- Stage 1 loads Stage 1.5
- Stage 1.5 loads the grub kernel and filesystem modules needed to locate Stage 2 (*/boot/grub2*)
- Stage 2 parses */boot/grub2/grub.cfg* and displays GRUB Menu
- When a (Linux) OS is selected, corresponding kernel image along with the appropriate initramfs images are loaded into memory.
- If Windows is selected, GRUB chainloads the Windows Bootloader.

The Chicken-and-Egg problem

The Chicken-and-Egg Problem

- The kernel needs to mount the root filesystem which can have one or more properties:
 - NFS
 - RAID
 - LVM
 - Encryption
- The necessary modules are stored in */lib/modules*
- In order to access them, the filesystem must be mounted.

The Chicken-and-Egg Problem

Solution?

Initial Ram Disk (initrd)

- A temporary filesystem that has all the stuff needed to mount the root filesystem.
- Initial Ram Disk (initrd) is an image of a block device.
- Grub makes it available as a special block device in memory (/dev/ram)
- The kernel mounts this device and uses it as it's root filesystem.
- Kernel must be compiled with the drivers needed to mount initrd.

Initrd vs Initramfs

initrd vs initramfs

- Initrd is made available as a block device that is mounted in memory.
- Initramfs is an archive that contains the directory structure of a typical filesystem.
- Archive is extracted to create a tmpfs (temporary filesystem)
- Mounted in memory to create initramfs
- Initramfs introduced in Linux Kernel v2.6.13

The Kernel

The Kernel

- Initializes and configures memory and various hardware.
- Extract initramfs into a tempfs and mount it in memory
- Mount the root filesystem as read-only.
- Start the first program.

SysV Init

SysV-Init

- The first program started by the kernel, thus has a PID of 1.
- A daemon that runs throughout the lifetime of the system until it shuts down. Manages all other daemons and programs.
- Program is located at “/sbin/init”

The rc.sysinit script

- Located at “/etc/rc.d/rc.sysinit”, this is the first script that is executed by init and it will perform the following tasks:
 - Set the system’s hostname
 - Unmount initramfs
 - Sets kernel parameters as defined in /etc/sysctl.conf.
 - Start devfs
 - Mount procfs and sysfs
 - Dumps the current contents of the kernel ring buffer into /var/log/dmesg
 - Process /etc/fstab (mounting and running fsck)
 - Enable RAID and LVM
- After the rc.sysinit script is executed, the relevant runlevel scripts are executed.

Runlevels

- A runlevel describes the state of a system with regards to the services and functionality that is available.
- There are a total of 7 runlevels which are defined as follows:
 - 0: Halt or shutdown the system
 - 1: Single user mode
 - 2: Multi-user mode, without networking
 - 3: Full multi user mode, with networking
 - 4: Officially not defined; Unused
 - 5: Full multi user with NFS and graphics (typical for desktops)
 - 6: Reboot
- Default runlevel is configured in the config file [/etc/inittab](#)

Runlevel Scripts

- Runlevel scripts are responsible for starting and stopping services that will provide the functionality for the respective runlevel.
- Located under “/etc/rc.d”:

```
$ ls -l /etc/rc.d
total 60
drwxr-xr-x 2 root root 4096 Aug 26 09:19 init.d
-rwxr-xr-x 1 root root 2617 Aug 17 2017 rc
drwxr-xr-x 2 root root 4096 Aug 26 09:19 rc0.d
drwxr-xr-x 2 root root 4096 Aug 26 09:19 rc1.d
drwxr-xr-x 2 root root 4096 Aug 26 09:19 rc2.d
drwxr-xr-x 2 root root 4096 Aug 26 09:19 rc3.d
drwxr-xr-x 2 root root 4096 Aug 26 09:19 rc4.d
drwxr-xr-x 2 root root 4096 Aug 26 09:19 rc5.d
drwxr-xr-x 2 root root 4096 Aug 26 09:19 rc6.d
-rwxr-xr-x 1 root root 220 Jul 2 06:56 rc.local
-rwxr-xr-x 1 root root 20108 Aug 17 2017 rc.sysinit
```

Runlevel Scripts

- Naming convention of the scripts is as follows:
- Starts with “S” or “K”. “S” scripts are used to start the service and “K” scripts are used to stop the service.
- Followed by a number that specifies the order in which the scripts are to be executed.
- Ends with the name of the service.

rc.local script

- After the runlevel scripts have been executed, the last script to run is the `/etc/rc.d/rc.local` script.
- You can add custom Bash commands you want to be executed at boot.
- The login prompt will show up after the `rc.local` script has been executed.

SysV-Init Important Commands

- `$ service httpd start`
- `$ service httpd status`
- `$ chkconfig httpd on`
- `$ runlevel`
- `$ telinit 3`
- `$ sed -i "s/^id:.*:initdefault:/id:3:initdefault/" /etc/inittab`

SystemD Boot Process

SystemD

- SystemD is a system and service manager that was designed to replace SysV-Init which has the following limitations:
- Services are started sequentially.
- Longer boot times.
- No easy and straightforward way to monitor running services.
- Dependencies have to be handled manually.

SystemD Units

- Every resource that is managed by SystemD is called a unit.
- A unit is a plain-text file that stores information about any one of the following:
 - a service
 - a socket
 - a device
 - a mount point,
 - an automount point
 - a swap file or partition
 - a start-up target
 - a watched file system path
 - a timer controlled and supervised by [systemd](#)
 - a resource management slice or a group of externally created processes.

SystemD Target Units

- Runlevels were replaced by “target” units,
- Target files are used to group together units that are needed for that specific target (services, sockets, devices, etc.)

SysV-Init Runlevel	SystemD Start-up Target
0: Halt or shutdown the system	poweroff.target
1: Single User mode	rescue.target
2: Multi-user mode, without networking	multi-user.target
3: Full multi user mode, with Networking	multi-user.target
4: Undefined	multi-user.target
5: Full multi-user mode with networking and graphical desktop.	graphical.target
6: Reboot	reboot.target

SystemD Service Units

- Runlevel Scripts were replaced by systemd unit files (mostly service unit files)
- A service unit file will define how that service is started and stopped.

```
[Unit]
Description=OpenSSH server daemon
Documentation=man:sshd(8) man:sshd_config(5)
After=network.target sshd-keygen.service
Wants=sshd-keygen.service

[Service]
Type=notify
EnvironmentFile=/etc/sysconfig/ssh
ExecStart=/usr/sbin/sshd -D $OPTIONS
ExecReload=/bin/kill -HUP $MAINPID
KillMode=process
Restart=on-failure
RestartSec=42s

[Install]
WantedBy=multi-user.target
```

```
[Unit]
Description=The Apache HTTP Server
Wants=httpd-init.service
After=network.target remote-fs.target nss-lookup.target httpd-init.service
Documentation=man:httpd.service(8)

[Service]
Type=notify
Environment=LANG=C

ExecStart=/usr/sbin/httpd $OPTIONS -DFOREGROUND
ExecReload=/usr/sbin/httpd $OPTIONS -k graceful
# Send SIGWINCH for graceful stop
KillSignal=SIGWINCH
KillMode=mixed
PrivateTmp=true

[Install]
WantedBy=multi-user.target
```


SystemD Dependencies

- There are 3 main ways of defining the dependencies of unit (target or service):
 - “**Wants=**” statements inside the unit files.
 - “**Requires=**” statements inside the unit files.
 - Special “**.wants**” directories associated with each target unit file found under the directory **/etc/systemd/system**
- When you enable a service using `systemctl`, `systemd` creates a symbolic link in the default target’s “**.wants**” directory that points to the service unit file.

SystemD Boot Process

- When a system boots into a specific target, all the dependent units are activated (services, mounts, sockets, etc.)
- All the tasks that are performed by the rc.sysinit script are replaced by “sysinit.target”.
- `basic.target`, `multi-user.target`, `graphical.target` and `rescue.target` all have `sysinit.target` as one of their dependencies.
- By default, rc.local script will not run by default unless if you make it executable:
 - `chmod +x /etc/rc.d/rc.local`

System Calls



System Calls

- Interface between user-level processes and hardware devices.
 - CPU, memory, disks etc.
- Make programming easier:
 - Let kernel take care of hardware-specific issues.
- Increase system security:
 - Let kernel check requested service via syscall.
- Provide portability:
 - Maintain interface but change functional implementation.

POSIX APIs

- API = Application Programmer Interface.
 - Function defn specifying how to obtain service.
 - By contrast, a system call is an explicit request to kernel made via a software interrupt.
- Standard C library (**libc**) contains wrapper routines that make system calls.
 - e.g., malloc, free are libc routines that use the brk system call.
- POSIX-compliant = having a standard set of APIs.
- Non-UNIX systems can be POSIX-compliant if they offer the required set of APIs.

Linux System Calls (1)

Invoked by executing **int \$0x80**.

- Programmed exception vector number 128.
- CPU switches to kernel mode & executes a kernel function.
- Calling process passes **syscall number** identifying system call in **eax** register (on Intel processors).
- Syscall handler responsible for:
 - Saving registers on kernel mode stack.
 - Invoking syscall service routine.
 - Exiting by calling **ret_from_sys_call()**.

Linux System Calls (2)

- System call dispatch table:
 - Associates syscall number with corresponding service routine.
 - Stored in **sys_call_table** array having up to **NR_syscall** entries (usually 256 maximum).
 - nth entry contains service routine address of syscall n.

Initializing System Calls

- **trap_init()** called during kernel initialization sets up the IDT (interrupt descriptor table) entry corresponding to vector 128:
 - **set_system_gate(0x80, &system_call);**
- A system gate descriptor is placed in the IDT, identifying address of **system_call** routine.
 - Does not disable maskable interrupts.
 - Sets the descriptor privilege level (DPL) to 3:
 - Allows User Mode processes to invoke exception handlers (i.e. syscall routines).

The `system_call()` Function

- Saves syscall number & CPU registers used by exception handler on the stack, except those automatically saved by control unit.
- Checks for valid system call.
- Invokes specific service routine associated with syscall number (contained in **`eax`**):
 - `call *sys_call_table(0, %eax, 4)`
- Return code of system call is stored in **`eax`**.

Parameter Passing

- On the 32-bit Intel 80x86:
 - 6 registers are used to store syscall parameters.
 - **eax** (syscall number).
 - **ebx, ecx, edx, esi, edi** store parameters to syscall service routine, identified by syscall number.

Wrapper Routines

- Kernel code (e.g., kernel threads) cannot use library routines.
- `_syscall0 ... _syscall5` macros define wrapper routines for system calls with up to 5 parameters.
- e.g., `_syscall3(int,write,int,fd,
const char *,buf,unsigned int,count)`

Example: “Hello, world!”

```
.data                                # section declaration

msg:
    .string "Hello, world!\n"        # our dear string
    len = . - msg                    # length of our dear string

.text                                # section declaration

    # we must export the entry point to the ELF linker or
.global _start                       # loader. They conventionally recognize _start as their
    # entry point. Use ld -e foo to override the default.

_start:

# write our string to stdout

    movl    $len,%edx                # third argument: message length
    movl    $msg,%ecx                # second argument: pointer to message to write
    movl    $1,%ebx                  # first argument: file handle (stdout)
    movl    $4,%eax                  # system call number (sys_write)
    int     $0x80                    # call kernel

# and exit

    movl    $0,%ebx                  # first argument: exit code
    movl    $1,%eax                  # system call number (sys_exit)
    int     $0x80                    # call kernel
```

Linux Files Relating to Syscalls

- Main files:
 - arch/i386/kernel/entry.S
 - System call and low-level fault handling routines.
 - include/asm-i386/unistd.h
 - System call numbers and macros.
 - kernel/sys.c
 - System call service routines.

arch/i386/kernel/entry.S

```
.data
ENTRY(sys_call_table)
.long SYMBOL_NAME(sys_ni_syscall) /* 0 - old "setup()" system
                                   call*/
    .long SYMBOL_NAME(sys_exit)
    .long SYMBOL_NAME(sys_fork)
    .long SYMBOL_NAME(sys_read)
    .long SYMBOL_NAME(sys_write)
```

- Add system calls by appending entry to sys_call_table:

```
.long SYMBOL_NAME(sys_my_system_call)
```

include/asm-i386/unistd.h

- Each system call needs a number in the system call table:
 - e.g., `#define __NR_write 4`
 - `#define __NR_my_system_call nnn`, where `nnn` is next free entry in system call table.


kernel/sys.c

- Service routine bodies are defined here:

- e.g., **asmlinkage retval**

```
    sys_my_system_call (parameters) {  
        body of service routine;  
        return retval;  
    }
```

Kernel Modules



Kernel Modules

- Modules can be compiled and dynamically linked into kernel address space.
 - Useful for device drivers that need not always be resident until needed.
 - Keeps core kernel “footprint” small.
 - Can be used to “extend” functionality of kernel too!

Example: “Hello, world!”

```
#define MODULE
#include <linux/module.h>
int init_module(void) {
    printk("<1>Hello, world!\n");
    return 0;
}
void cleanup_module(void) {
    printk("<1>Goodbye cruel world ☹️\n");
}
```

Using Modules

- Module object file is installed in running kernel using **insmod module_name**.
 - Loads module into kernel address space and links unresolved symbols in module to symbol table of running kernel.

The Kernel Symbol Table

- Symbols accessible to kernel-loadable modules appear in **/proc/ksyms**.
 - **register_syntab** registers a symbol table in the kernel's main table.
 - Real hackers export symbols from the kernel by modifying **kernel/ksyms.c** 😊