

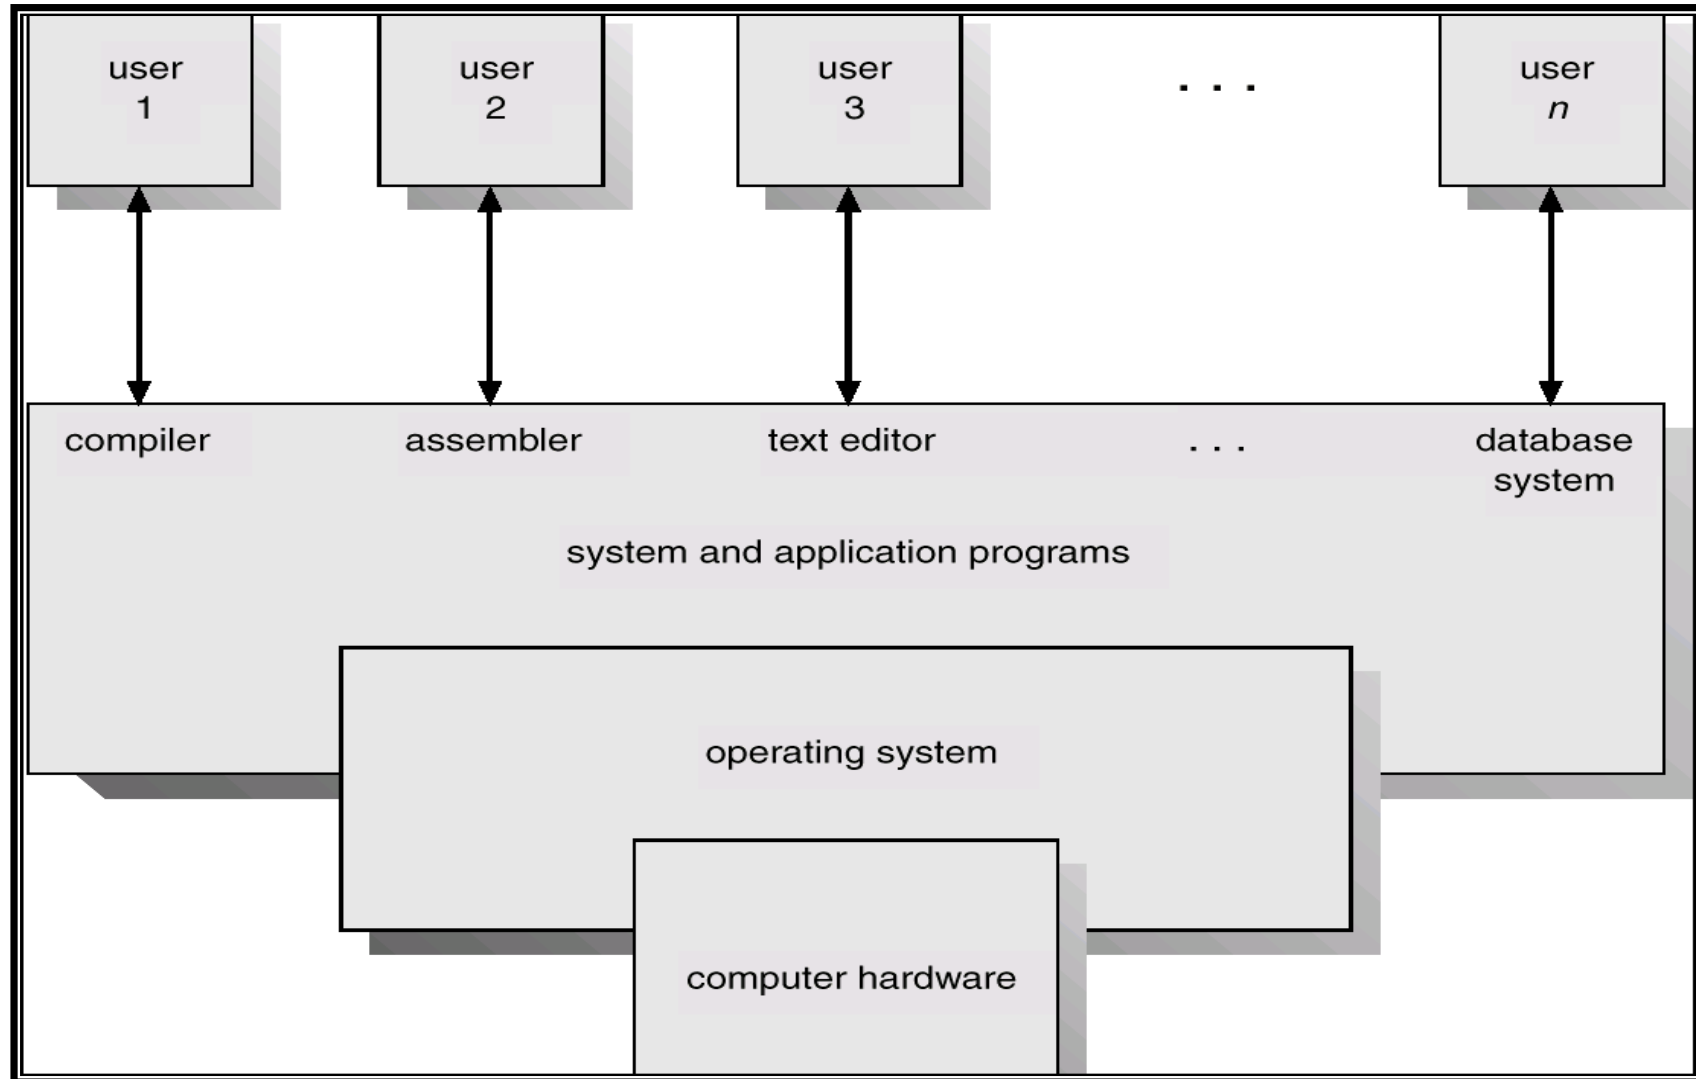
Advanced OS - Kernels

Manish Shrivastava

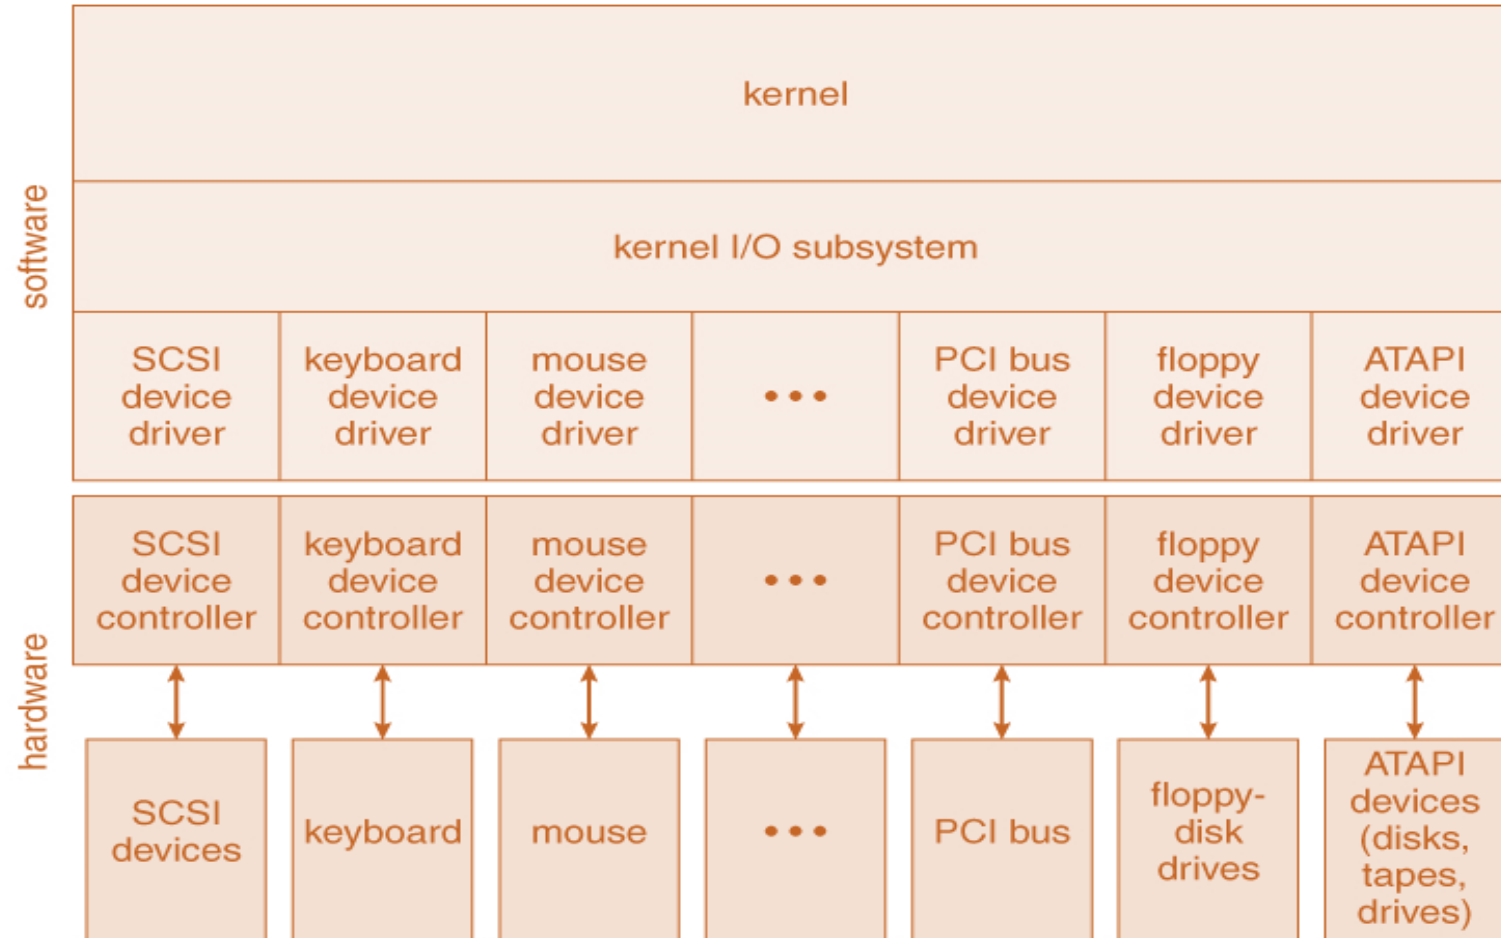
Computer OS components

- **Hardware** – provides basic computing resources (CPU, memory, I/O devices).
- **Operating system** – controls and coordinates the use of the hardware among the various application programs for the various users.
- **Applications programs** – define the ways in which the system resources are used to solve the computing problems of the users (compilers, database systems, video games, business programs).
- **Users** (people, machines, other computers).

Computer OS components



Kernel




Types of Kernels

1. Monolithic Kernel: Kernels where the user services and the kernel services are implemented in the same memory space
2. Microkernel: the user services and kernel services are implemented into different spaces
3. Hybrid Kernel: Duh!!

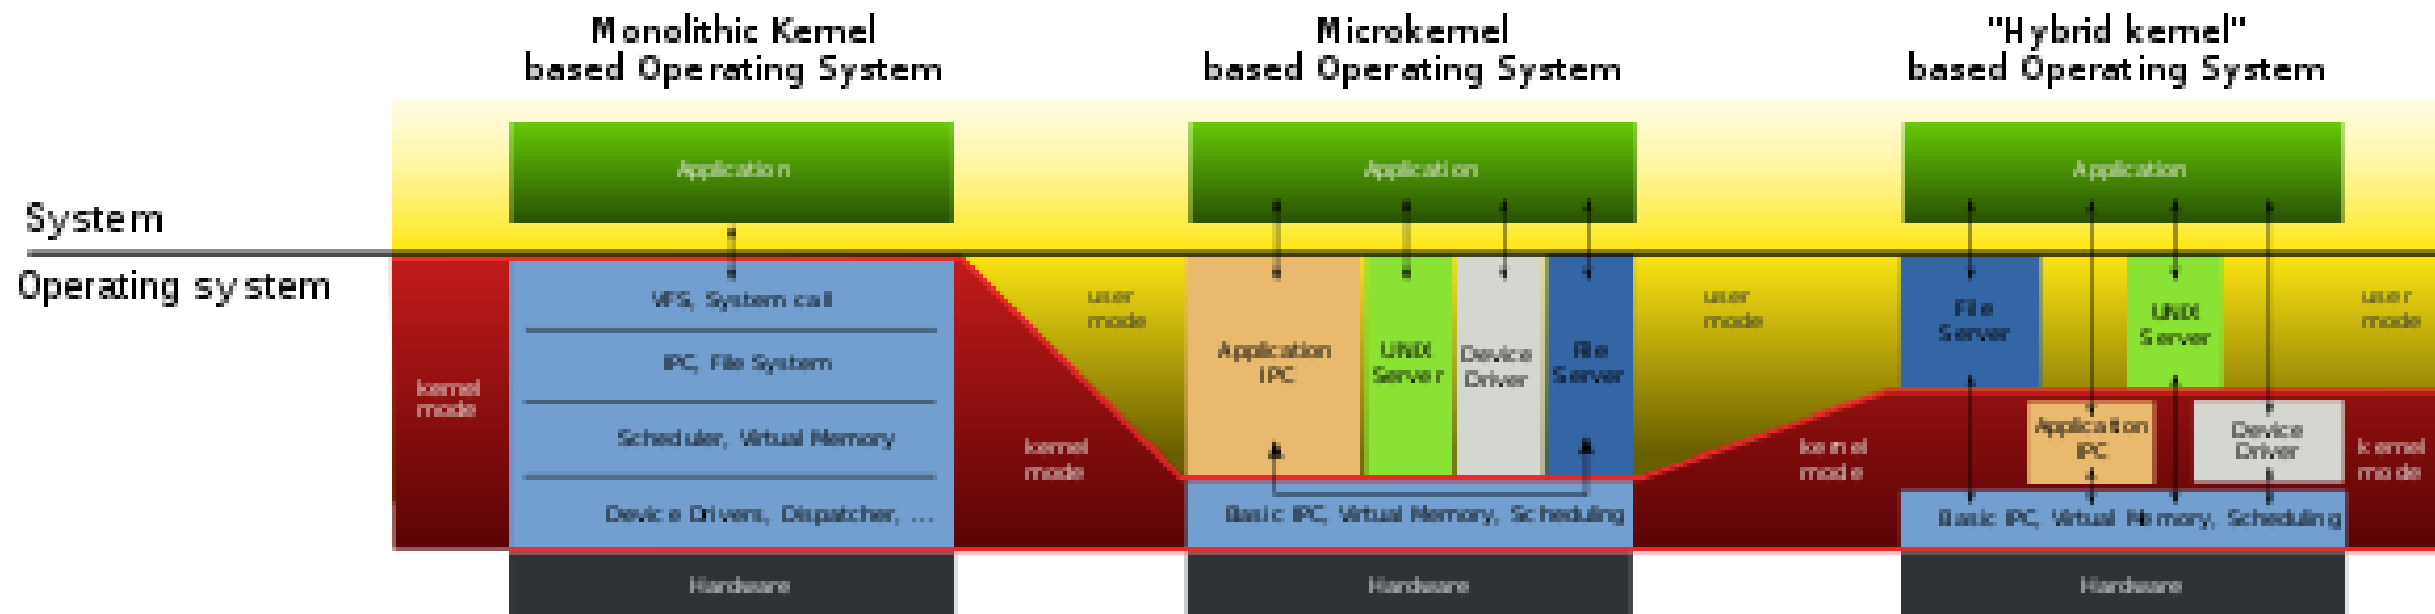
Early structure: Monolithic

- Traditionally, OS's (like UNIX) were built as a **monolithic** entity:

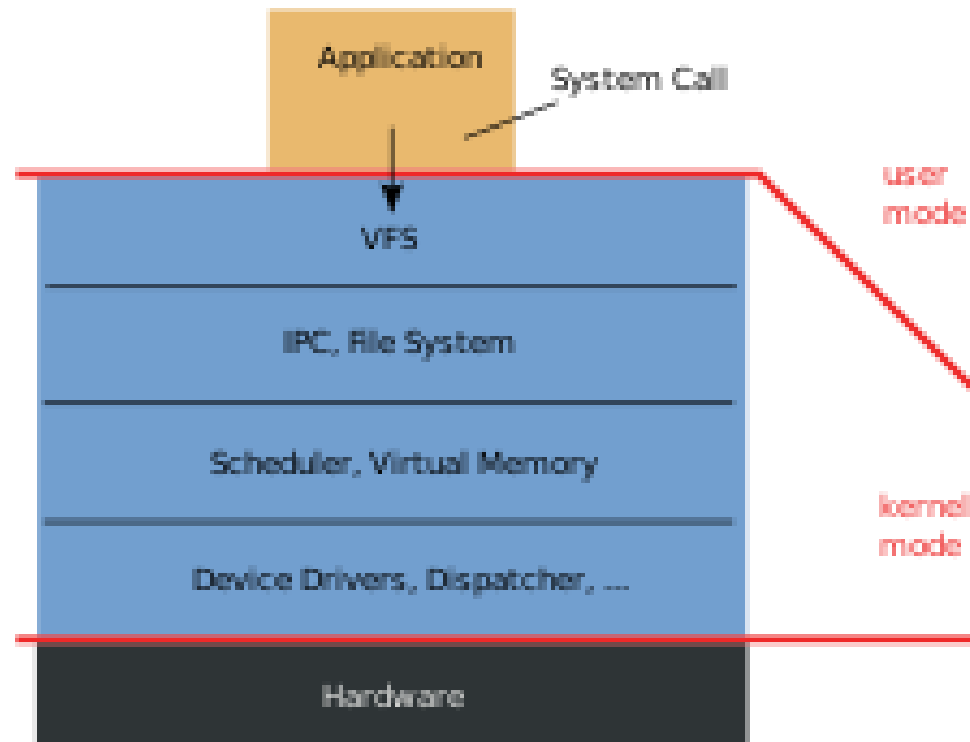


Monolithic design


- Major advantage:
 - cost of module interactions is low (procedure call)
- Disadvantages:
 - hard to understand
 - hard to modify
 - unreliable (no isolation between system modules)
 - hard to maintain
- What is the alternative?
 - find a way to organize the OS in order to simplify its design and implementation



Monolithic Kernel based Operating System



Microkernel based Operating System



Layering

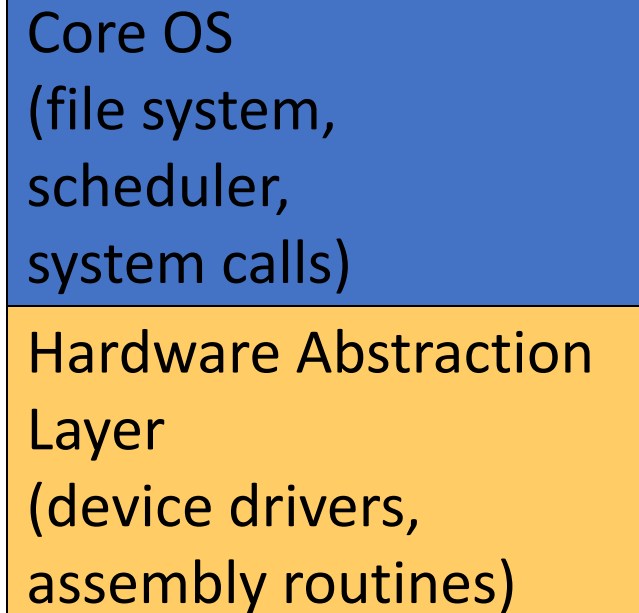
- The traditional approach is layering
 - implement OS as a set of layers
 - each layer presents an enhanced 'virtual machine' to the layer above
- The first description of this approach was Dijkstra's THE system
 - Layer 5: **Job Managers**
 - Execute users' programs
 - Layer 4: **Device Managers**
 - Handle devices and provide buffering
 - Layer 3: **Console Manager**
 - Implements virtual consoles
 - Layer 2: **Page Manager**
 - Implements virtual memories for each process
 - Layer 1: **Kernel**
 - Implements a virtual processor for each process
 - Layer 0: **Hardware**
- Each layer can be tested and verified independently

Problems with layering

- Imposes hierarchical structure
 - but real systems are more complex:
 - file system requires VM services (buffers)
 - VM would like to use files for its backing store
 - strict layering isn't flexible enough
- Poor performance
 - each layer crossing has **overhead** associated with it
- Disjunction between model and reality
 - systems modeled as layers, but not really built that way

Hardware Abstraction Layer

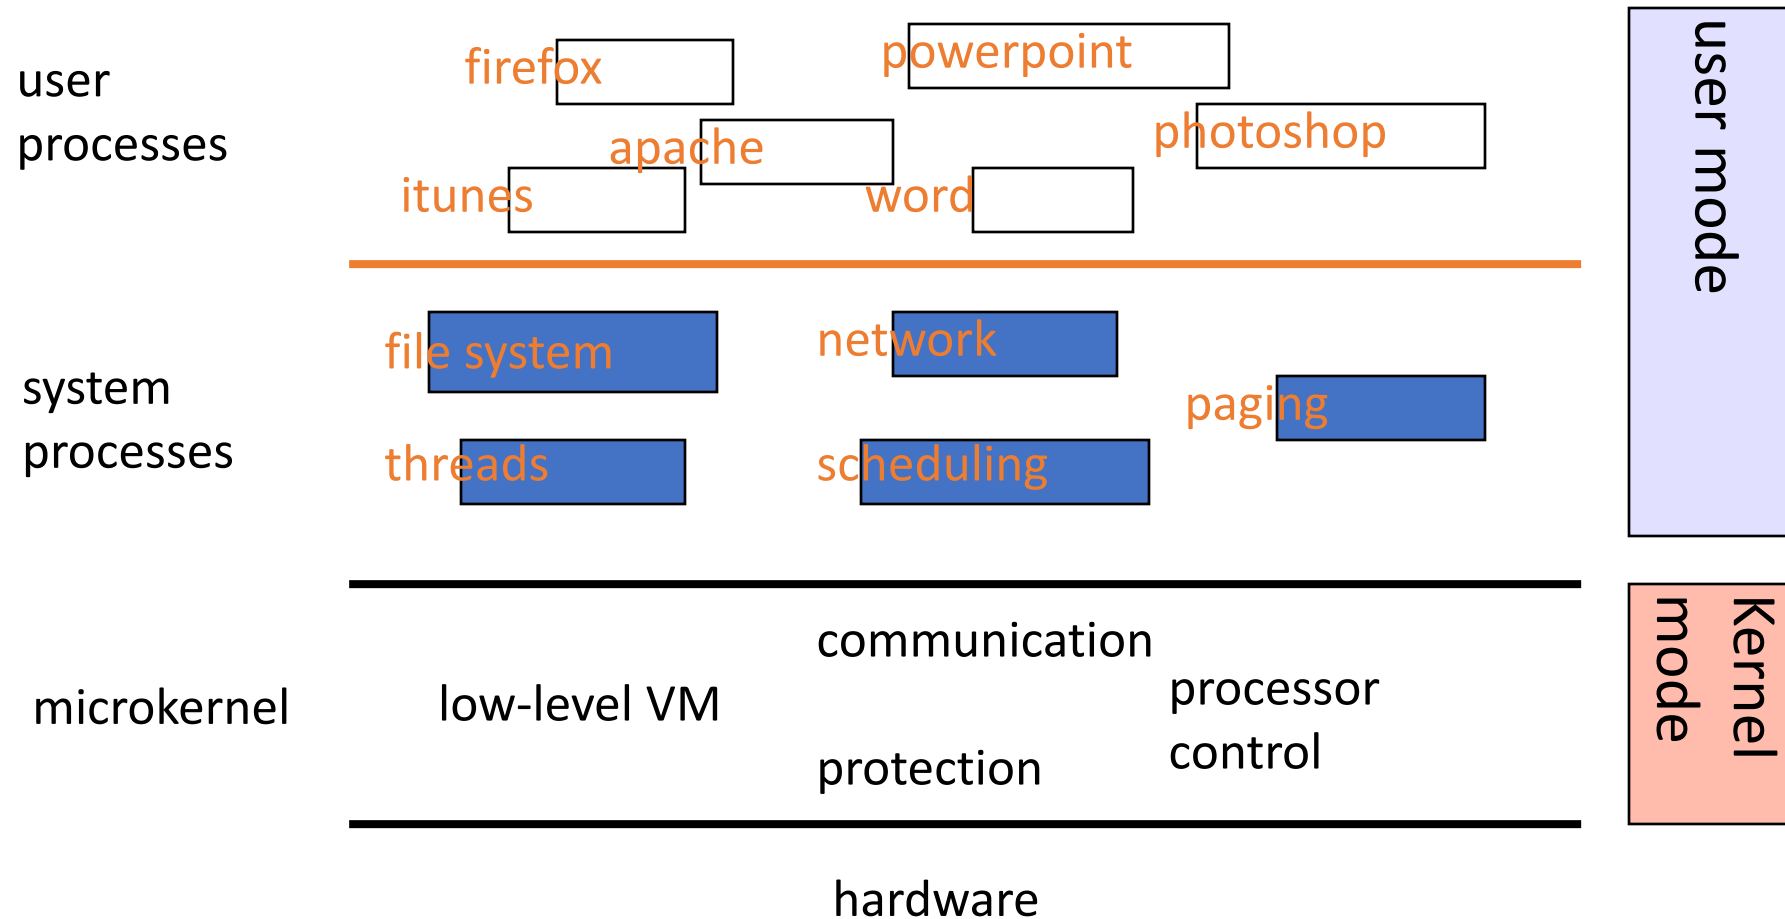
- An example of layering in modern operating systems
- Goal: separates hardware-specific routines from the “core” OS
 - Provides portability
 - Improves readability



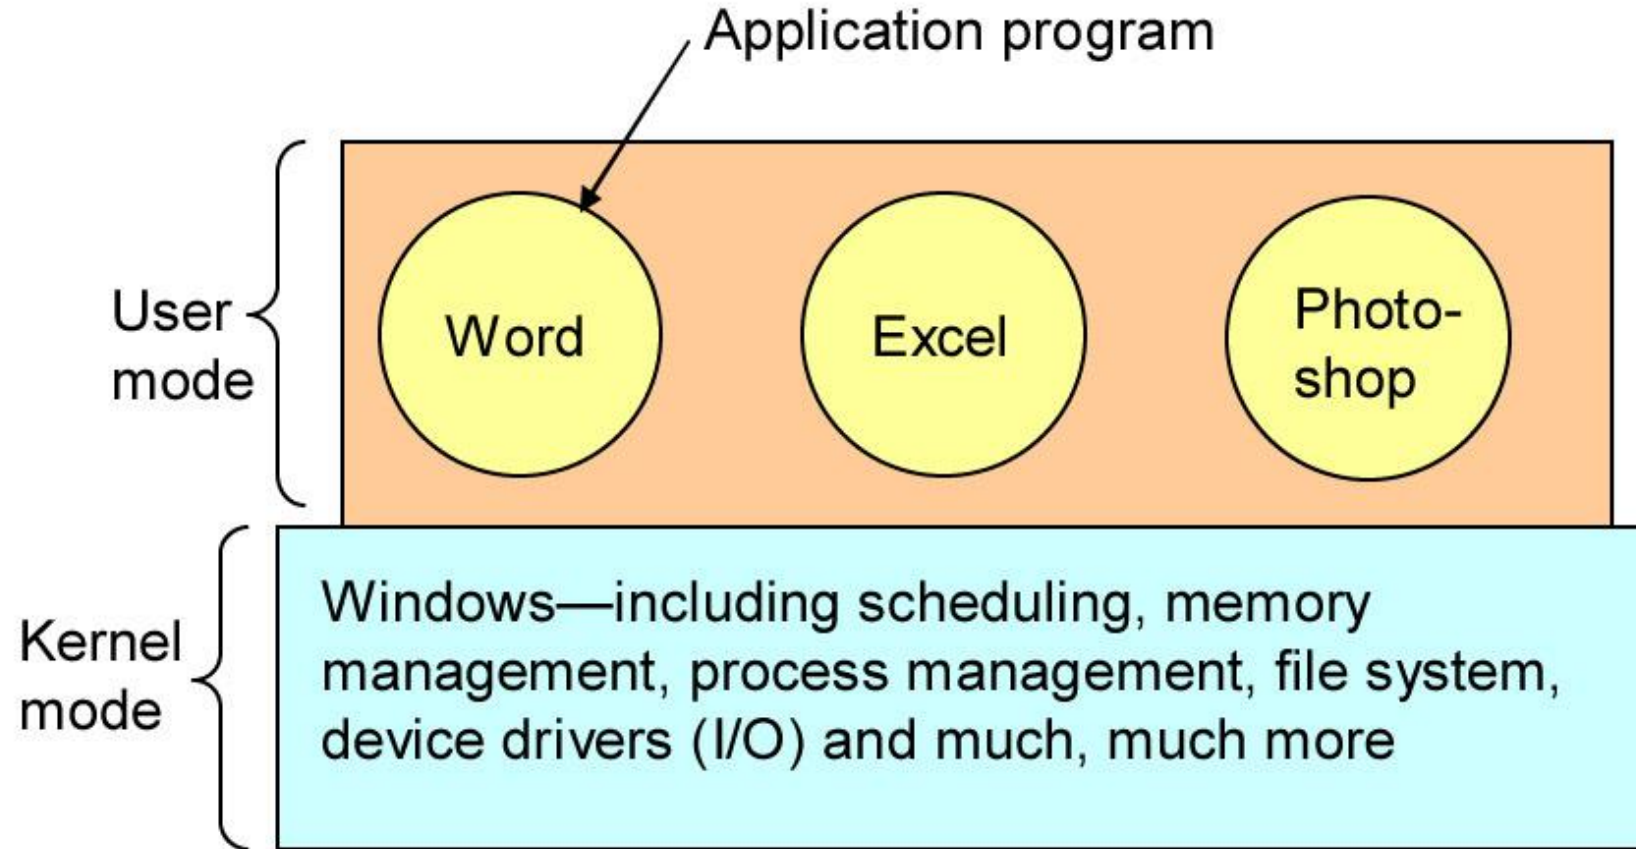
Microkernels

- Popular in the late 80's, early 90's
 - recent resurgence of popularity
- Goal:
 - minimize what goes in kernel
 - organize rest of OS as user-level processes
- This results in:
 - better reliability (isolation between components)
 - ease of extension and customization
 - poor performance (user/kernel boundary crossings)
- First microkernel system was Hydra (CMU, 1970)
 - Follow-ons: Mach (CMU), Chorus (French UNIX-like OS), OS X (Apple), in some ways NT (Microsoft)

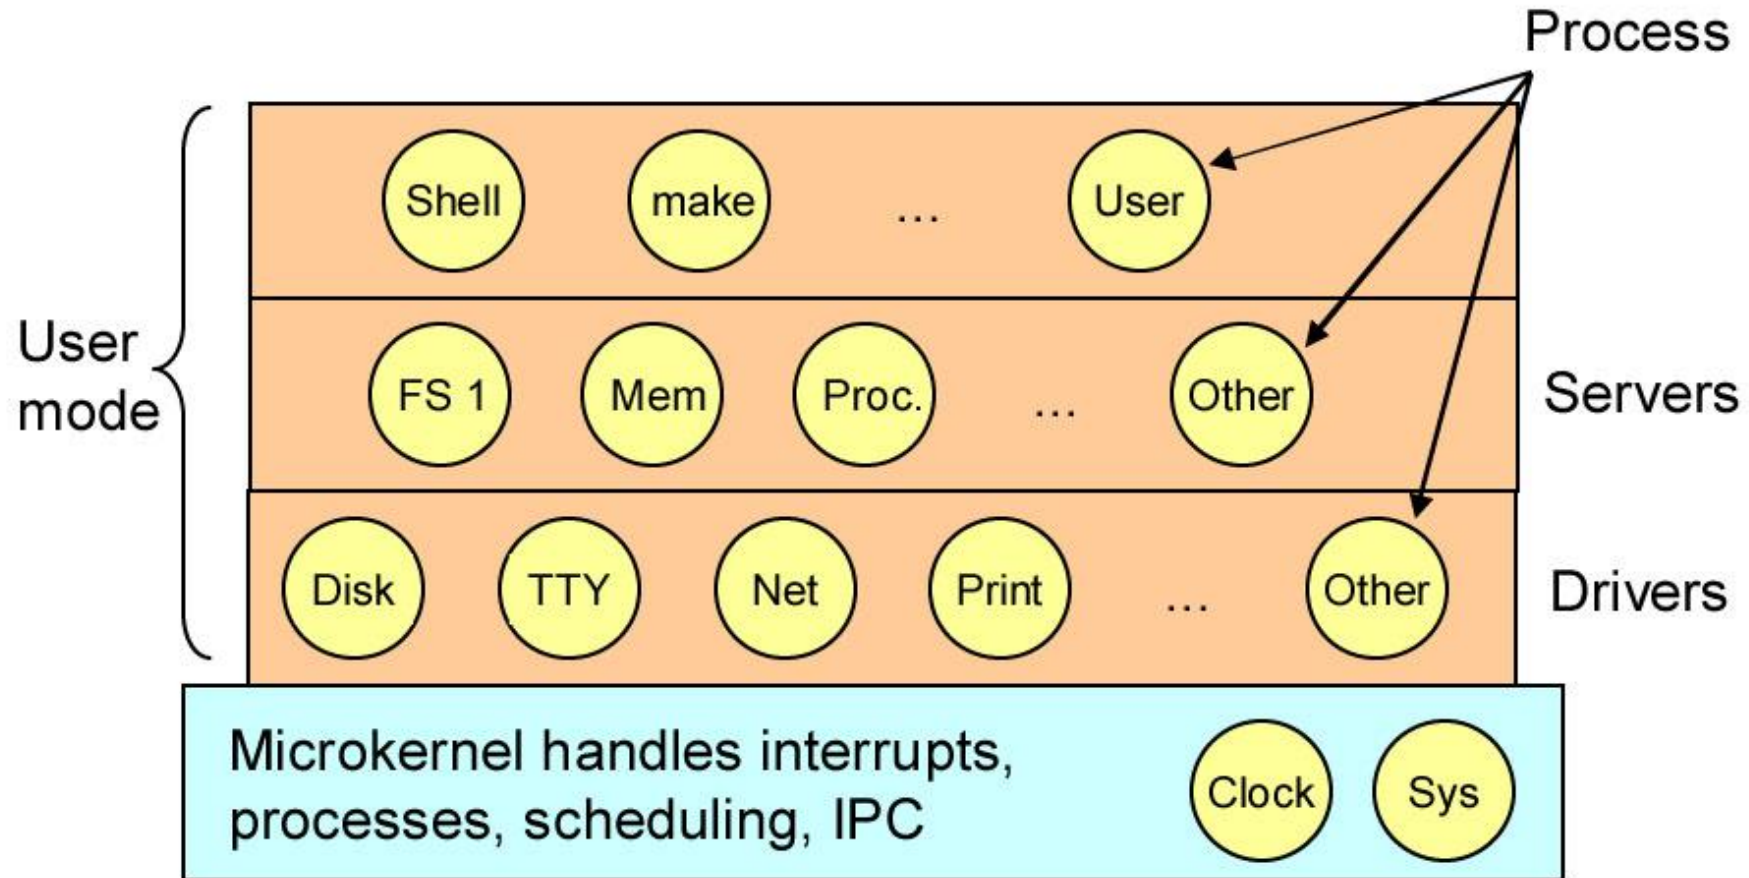
Microkernel structure illustrated



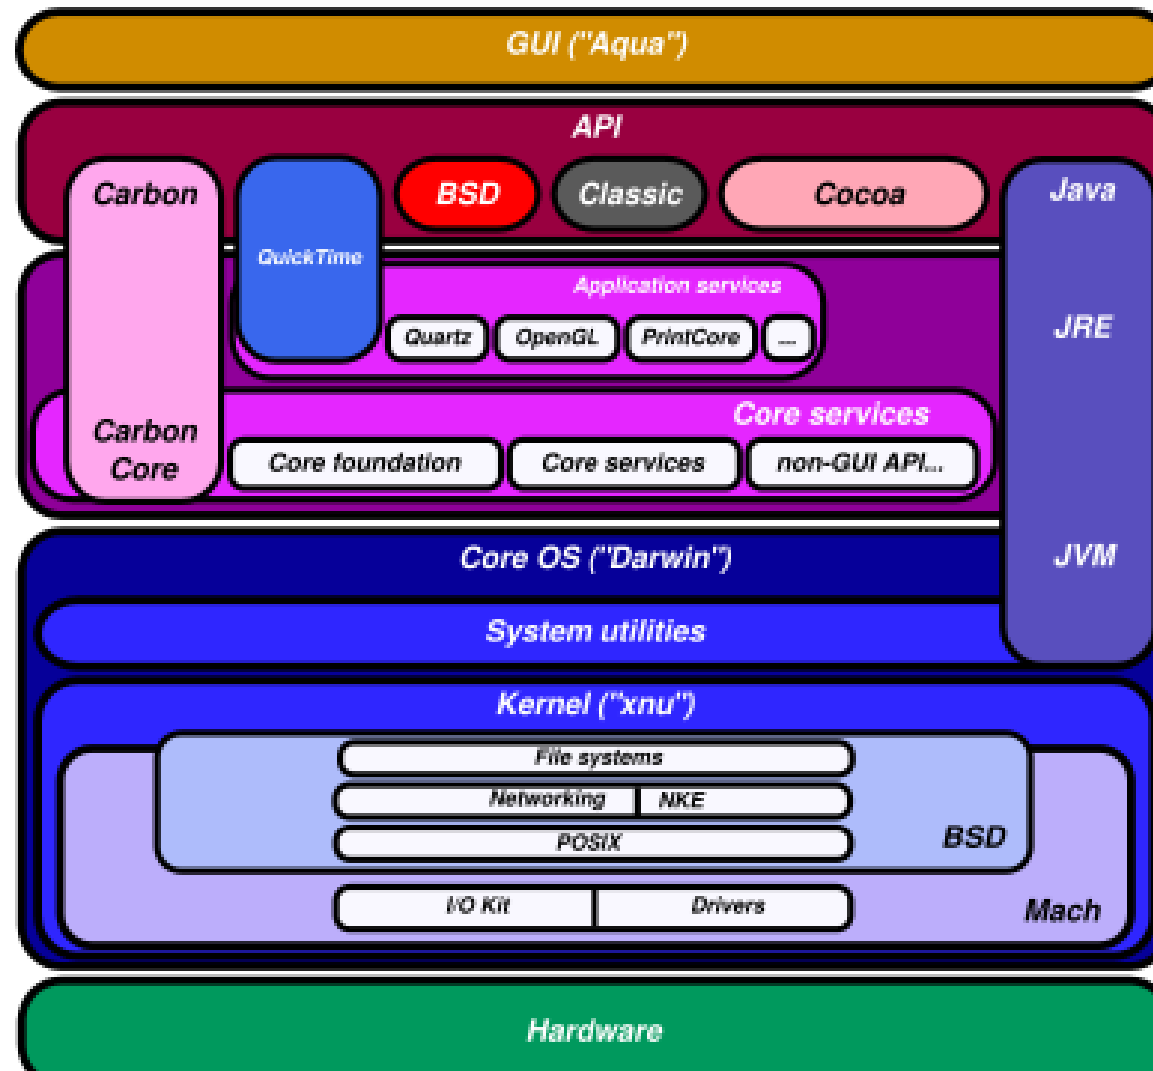
EXAMPLE: WINDOWS



ARCHITECTURE OF MINIX 3



Mac OS X



The Linux Kernel

Manish Shrivastava

Monolithic?

- What is a monolithic program?
- What is a Monolithic Kernel?

Process?

- What is a process?
- Is kernel a process?

Process?

- What is a process?
- Is kernel a process?
- No!
- Init is just the first process; it does not manage any processes or threads. It does create some, using the kernel syscalls `fork()` and `exec`.
- It is a Process Manager without being a process

Kernel vs. Application Coding

Two major differences:

- The core kernel has no standard libraries
- The core kernel must be a monolithic, statically linked library

No standard libraries:

- No libc (malloc, pthreads, string handling, etc.)
- Partly because of chicken/egg situation
- Also, standard libraries can be too slow

No dynamic loading of shared libraries for the core kernel

Two Big Problems

1. A totally static kernel would be enormous
 - About 20 million lines of code in 2015
 - Most of this is hardware drivers that are never used on any given platform
2. Programmers rely on “library” functions for efficiency and correctness
 - E.g., things like data structures, sleeping routines, etc.
 - What libraries to use? (hint: not glibc)
 - Kernel code has to be entirely self-contained
 - Also, chicken+egg problem, as glibc functions may in turn invoke the kernel via system calls

First Solution:

Loadable Kernel Modules

Kernel modules are kernel code that can be loaded dynamically:

- Can be loaded/unloaded whenever
- Runs in kernel mode
- Can access *exported* kernel variables and functions
- Can export variables and functions to kernel or other modules

Kernel and Kernel Modules

- Often, if you want to add something to the kernel you need to rebuild the kernel and reboot.
 - A “loadable kernel module” (LKM) is an object file that extends the base kernel.
 - Exist in most OSes
 - Including Windows, FreeBSD, Mac OS X, etc.
 - Modules get added and removed as needed
 - To save memory, add functionality, etc.

Why Kernel Modules

Linux is a monolithic kernel. All functionality is compiled into the same static binary that is loaded into memory on boot.

Without modules, ***the entire kernel*** would need to be loaded into memory to boot a node. Problems?

- Waste of memory (embedded systems)
- Slower boot time
- Larger trusted computing base (TCB), more room for bugs

What are Kernel Modules used for?

- What pieces of code do we think might not be needed on every system that the kernel boots on?
 1. Device Drivers
 2. Architecture-specific code
- Lines of code (just `.c` files) for:
 1. Device Drivers: 1,583,159
 2. Everything else: 1,003,777
 - (Includes all of the architectures that we're not using!)

What are Kernel Modules used for?

- Beyond space savings, what else might modules be useful for?
 1. “Out-of-tree” functionality that is not accepted into “mainline” kernel
 2. Allowing users to load custom functionality at runtime
 3. Configuring/patching a running system without requiring a reboot

Modular?

- What does it mean to be modular (exactly)?
- How is that beneficial?
- How would a Modular kernel differ from a Monolithic one?

Linux Kernel Modules

- In general must be licensed under a free license.
 - Doing otherwise will taint the whole kernel.
 - A tainted kernel sees little support.
 - Might be a copyright problem if you redistribute.
- The Linux kernel changes pretty rapidly, including APIs etc.
 - This can make it a real chore to keep LKMs up to date.
 - Also makes a tutorial a bit of a pain.

Module Implementation

Must define:

- An initialization function called on load
- An exit function called on unload

The init function must be self contained!

- Must unwind actions if initialization cannot complete successfully
- E.g. if you `kmalloc()` space but don't free it, that physical memory is now lost (until system restart)

You can also pass parameters to modules at load time.

Creating a module

- All modules need to define functions that are to be run when:
 - The module is loaded into the kernel
 - The module is removed from the kernel
- We just write C code (see next slide)
- We need to compile it as a kernel module.
 - We invoke the kernel's makefile.
 - `sudo make -C /lib/modules/xxx/build M=$PWD modules`
 - This makes (as root) using the makefile in the path specified.
 - I think it makes all C files in the directory you started in
 - Creates .ko (rather than .o) file
 - Xxx is some kernel version/directory

Simple module

```
#include <linux/init.h>
#include <linux/module.h>
#include <linux/kernel.h>

MODULE_LICENSE("Dual BSD/GPL");

static int hello_init(void) {
    printk("<1> Hello World!\n");
    return 0;
}

static void hello_exit(void) {
    printk("<1> Bye world!\n");
}

module_init(hello_init);
module_exit(hello_exit);
```

- MODULE_LICENSE
 - Required.
 - Short list of allowed licenses.
- Printk()
 - Kernel print.
 - Prints message to console and to log.
 - <1> indicates high priority message, so it gets logged.
- Module_init()
 - Tells system what module to call when we first load the module.
 - TIMTOWTDI
- Module_exit()
 - Same but called when module released.

Modules:

Listing, loading and removing

- From the command line:
 - lsmod
 - List modules.
 - insmod
 - Insert module into kernel
 - Adds to list of available modules
 - Causes function specified by `module_init()` to be called.
 - rmmod
 - Removes module from kernel

lsmod

Module	Size	Used by
memory	10888	0
hello	9600	0
binfmt_misc	18572	1
bridge	63776	0
stp	11140	1 bridge
bnep	22912	2
video	29844	0

insmod

- Very (very) simple
 - **insmod `xxxxx.ko`**
 - Says to insert the module into the kernel

Other (better) way to load a module

- **Modprobe** is a smarter version of insmod.
 - Actually it's a smarter version of insmod, lsmod and rmmod...
 - It can use short names/aliases for modules
 - It will first install any dependent modules
- We'll use insmod for the most part
 - But be aware of modprobe

So?

```
#include <linux/init.h>
#include <linux/module.h>
#include <linux/kernel.h>

MODULE_LICENSE("Dual BSD/GPL");

static int hello_init(void) {
    printk("<1> Hello World!\n");
    return 0;
}

static void hello_exit(void) {
    printk("<1> Bye world!\n");
}

module_init(hello_init);
module_exit(hello_exit);
```

- When insmod, log file gets a “Hello World!”
- When rmmod, that message prints to log (and console...)
- It’s not the name, it’s the module_init().

Modules?

- There are a number of different reasons one might have a module
 - But the main one is to create a device driver
 - It's not realistic for Linux to have a device driver for all possible hardware in memory all at once.
 - Would be too much code, requiring too much memory.
 - So we have devices as modules
 - Loaded as needed.

Device driver

(Thanks Wikipedia!)

- A device driver is a computer program allowing higher-level computer programs to interact with a hardware device.
 - A driver typically communicates with the device through the computer bus or communications subsystem to which the hardware connects.
 - When a calling program invokes a routine in the driver, the driver issues commands to the device.
 - Drivers are hardware-dependent and operating-system-specific.

Devices in Linux (1/2)

- There are special files called “device files” in Linux.
 - A user can interact with it much like a normal file.
 - But they *generally* provide access to a physical device.
 - They are generally found in /dev and /sys
 - /dev/fb is the frame buffer
 - /dev/ttyS0 is one of the serial ports

```
crw-rw---- 1 root dialout
```

- Not all devices files correspond to physical devices.

- Pseudo-devices.
 - Provide various functions to the programmer
 - /dev/null
 - Accepts and discards all input; produces no output.
 - /dev/zero
 - Produces a continuous stream of NULL (zero value) bytes.

```
4, 64 Jun 20 13:01 ttyS0
```

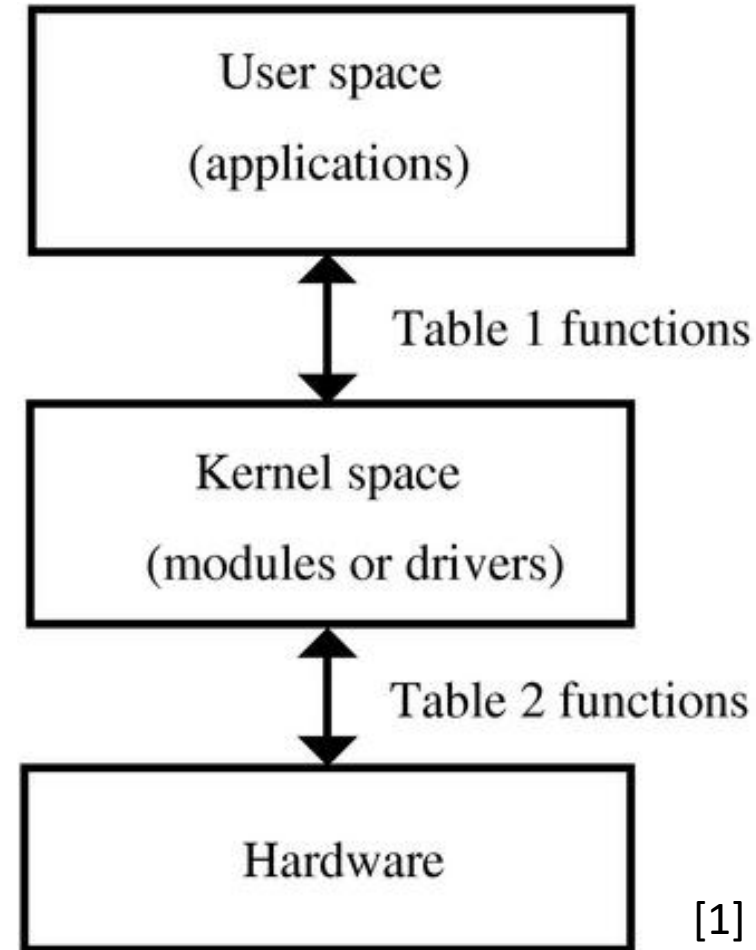
etc.

Devices in Linux (2/2)

- Pretty clearly you need a way to connect the device file to the actual device
 - Or pseudo device for that matter
- We want to be able to “fake” this by writing functions that handle the file I/O.
 - So we need to associate functions with all the things we can do with a file.
 - Open, close.
 - Read, write.
- Today we’ll talk about all that...

Kernel vs. User space

- User Space
 - End-user programs. They use the kernel to interface to the hardware.
- Kernel Space
 - Provides a standard (and hopefully multi-user secure) method of using and sharing the hardware.
 - Private function member might be a good analogy.



What is a “device”?

- Linux devices are accessed from user space in exactly the same way files are accessed.
 - They are generally found in /dev and /sys
- To link normal files with a kernel module, each device has a “major number”
 - Each device also has a “minor number” which can be used by the device to distinguish what job it is doing.

```
% ls -l /dev/fd0 /dev/fd0u1680
brwxrwxrwx    1 root   floppy    2,   0 Jul  5  2000 /dev/fd0
brw-rw----    1 root   floppy    2,  44 Jul  5  2000
/dev/fd0u1680
```

Two floppy devices. They are actually both the same bit of hardware using the same driver (major number is 2), but one is 1.68MB the other 1.44.

Creating a device

- **`mknod /dev/memory c 60 0`**
 - Creates a device named `/dev/memory`
 - Major number 60
 - Minor number 0
- Minor numbers are passed to the driver to distinguish different hardware with the same driver.
 - Or, potentially, the same hardware with different parameters (as the floppy example)

Second Solution: Kernel “Libraries”

Kernel libraries re-implement a lot of the functionality programmers expect in user space

- Are statically compiled into the kernel
- Automatically available just by including relevant header
- Built to be kernel-safe (sleeping, waiting, locking, etc. is done properly)

Features:

- Utilities: kmalloc, kthreads, string parsing, etc.
- Containers: hash tables, binary trees etc.
- Algorithms: sorting, compression

Mostly found under `/lib` **and** `/include/linux`

Kernel “Libraries”

Mostly found under `/lib` **and** `/include/linux`

Many kernel “libraries” have clear analogues to user space libraries:

- malloc/free vs kmalloc/kfree
- pthread_create vs kthread_create
- sleep/usleep/nanosleep (user) vs msleep (kernel)

Some, however, are a bit different:

- e.g., linked list implementation

Example: Linked Lists

See also:

- `/include/linux/list.h`
- `/include/linux/types.h`

Moral of the story:

Always search for functionality before writing it yourself.

Notes:

- One important note is that module stuff is written in kernel space.
 - That means you can't do a lot of things you might want to!
 - File I/O is a really bad idea
 - See next slide.
 - Talking to memory-mapped I/O devices requires effort
 - Still have virtual memory
 - Things like malloc don't quite work
 - Thus kmalloc, kprint, etc.
 - Can be an unpleasant place to live...